

Supervised learning. Decision tree learning (DTs)

SUMMARY

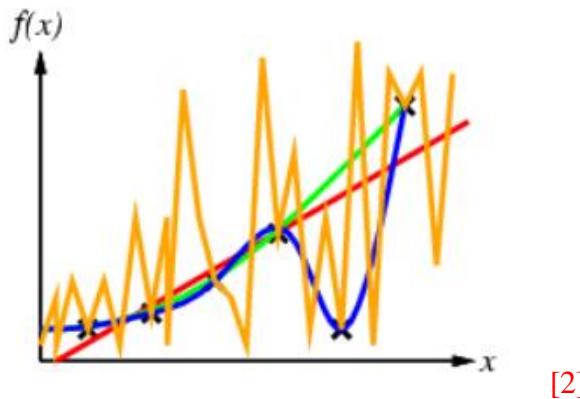
1. Issues in ML.....	1
2. Supervised learning.....	1
3. Eager inductive learners.....	3
4. Decision tree (DT) learning	13
5. DT related research topics.....	17

1. Issues in ML

- what algorithms can approximate functions well (and when)?
- how does the number of training examples influence performance?
- how does complexity of hypothesis representation impact learning performance?
- how does noisy data influence performance?
- what are the theoretical limits of learnability?
- ...

2. Supervised learning

- example: Inductive learning
- learn a function from examples
- Occam's razor
 - prefer the simplest hypothesis consistent with data
 - complex models tend not to generalize well



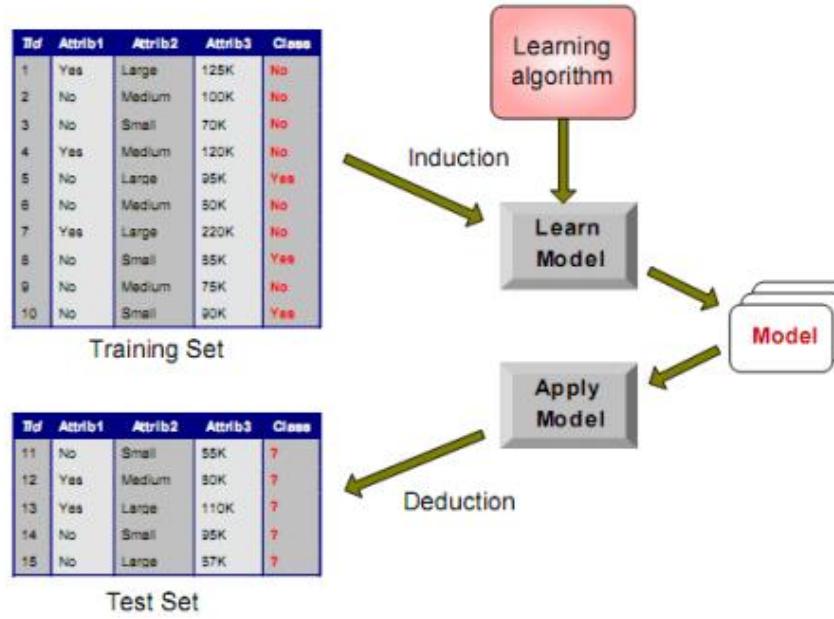
[2]

- **Classification vs. regression**
 1. learning a relationship between an input (vector x) and an output (y) from data
 2. **classification** estimates the discrete output y (known as the class)
 - types
 - One Class Classification (OCC)
 - Imbalanced data
 - Anomaly/outlier detection, open set recognition (OSR)
 - E.g., OSVM,...
 - Binary classification
 - Multi-class classification
 - predicting tumor cells as benign or malignant
 - categorizing news according to their domain (finance, sports, weather, etc)
 - classifying secondary structures of proteins
 - classifying credit card transactions
 -
 - 3. **regression** estimates the function f such that $y=f(x)$ with some confidence measure
 - predicting the exchange rate
 - predicting the price of houses
 - predicting the age of a person
 -
- **Offline vs. online** learning
 1. **Offline (batch) learning**
 - generates the best predictor by learning on the entire data set at once
 2. **Online learning**
 - data becomes available in a sequential order and is used to update the best predictor for future data at each step
 - learning incrementally (one instance at time)
 - useful where is computationally unfeasible to train over the entire data set (e.g. stock price prediction)
 - time dependent
- **Eager vs. lazy** learning
 1. **Eager learning**
 - an example of offline learning
 - the system tries to construct a general, input independent target function during the training of the system
 - a global model (hypothesis) is built during the training step
 - slow training, fast evaluation
 - builds a global estimate of the target function
 - deal much better with noise
 - generally unable to provide good local approximations of the target function
 - e.g. ANNs, SVMs, DTs, NBC, etc
 2. **Lazy learning**
 - little or no offline processing of the training data
 - motivation
 - online learning
 - data set is continuously updated with new entries

- generalization beyond the training set is delayed until a query is made to the system
- simply stores the training examples
 - not really a training phase
 - fast “training”, slow evaluation
- e.g. instance-based learning methods (kNN, LWR, CBR)
 - e.g., kNN used in [online recommendation systems](#) (movies, music, items to sale, a.s.o)
- provide a local approximation for the target function, for each new query instance

3. Eager inductive learners

- Most of the inductive learning models are eager ones.
- Building an eager inductive learner



[2]

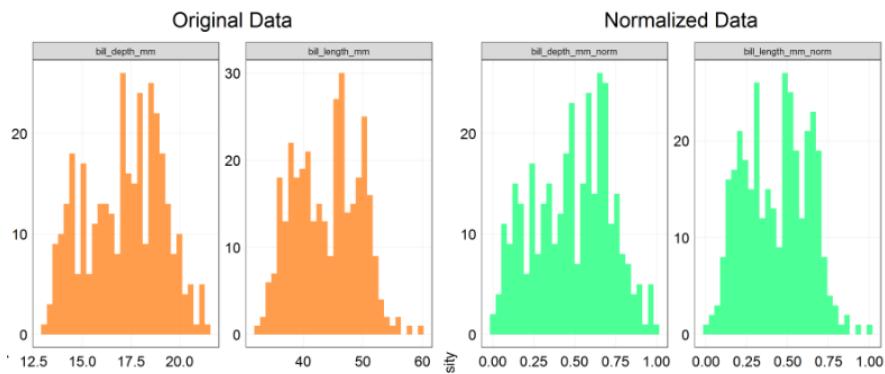
- train-validation-test sets
- e.g., 60-20-20 split

Data pre-processing and analysis

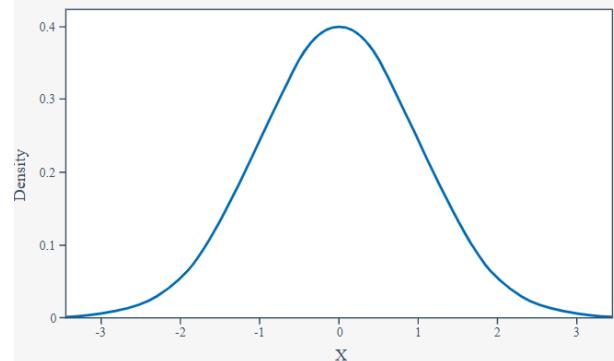
- data cleaning (cleansing)
- **normalization/standardization** (for numerical data), if needed
 - if X is a feature characterizing the data, $X = (X_1, X_2, \dots, X_n)$
 - normalization and standardization will not change the distribution of the data
 - **(min-max) normalization**

- changes the range of the data (i.e., the relative distances between individual data points)
 - does not alter the data distribution
- scale the data between 0 and 1

$$X'_i = \frac{X_i - \min}{\max - \min}$$
 - where \min and \max represent the *minimum* and *maximum* values of X
 - useful in ML algorithms that do not assume any distribution of data (kNN, neural networks).

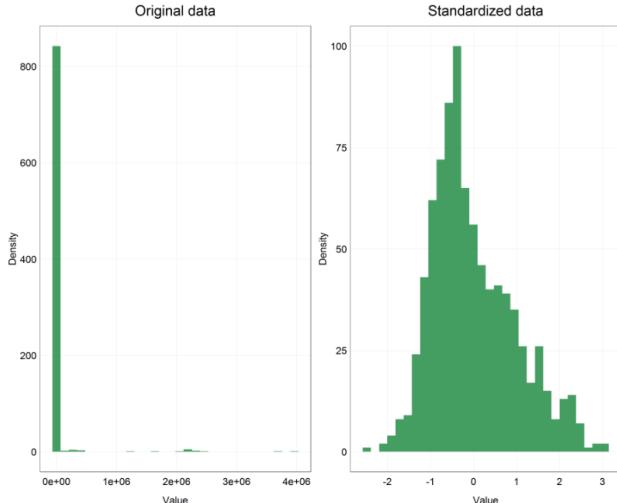


- **standard deviation-based normalization (standardization)**
 - standard normal distribution $N(0,1)$
 - 0 mean and 1 stdev (1 variance)



- after standardization the data will have zero mean and unit variance

$$X'_i = \frac{X_i - \mu}{\sigma}$$
 - μ represents the mean and σ represent the standard deviation of X
- preserves the general data distribution
 - shifts and scales a distribution
 - it does not change the shape of the distribution



- the transformed data are now centered on 0

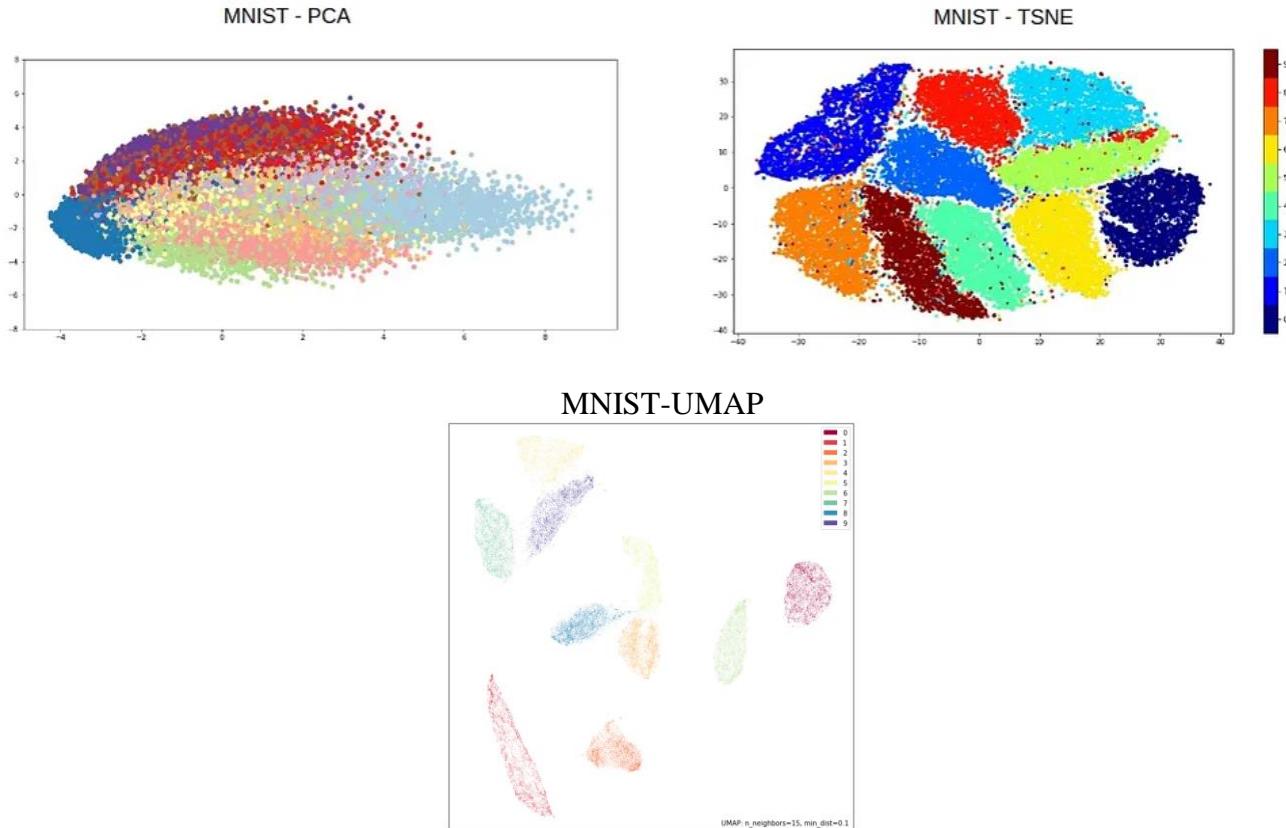
▪ **feature selection**

- **select relevant features**
 - correlation
 - independence tests
 - Pearson's χ^2 test
 - unsupervised learning
 - univariate filter methods
 - information theory, spectral analysis (manifold learning)
 - **univariate feature selection**
 - feature importance – statistical tests
 - select the most important features in a dataset
 - evaluate a feature's relationship with the target variable and select the ones that have the strongest correlation
 - *scikit-learn*
 - F-test for feature scoring
 - algorithms: Relief-based
 - *scikit-learn*

▪ **component reduction**

- build new input variables in order to reduce their number
- no loss of information about data distribution
- data vizualisation tools
- **Principal Component Analysis (PCA)**
 - map data into a space of lower dimensionality
 - linear mapping
 - preserving as much of the data's variance
- **Curvilinear Component Analysis (CCA)**
 - non-linear extension of PCA
 - is a kind of self-organizing map
- **t-Distributed Stochastic Neighbor Embedding (t-SNE)**
 - non-linear dimensionality reduction

- Uniform Manifold Approximation and Projection for Dimension Reduction ([UMAP](#))



- **others**

- outliers, isolated instances
- **imbalanced data**
 - data augmentation
 - oversampling
 - SMOTE (Synthetic Minority Oversampling TEchnique)
 - Cluster-based oversampling (CBOS)
 - Others
 - generative oversampling
 - [GANs](#), [VAE-based](#)
 - undersampling
 - algorithms/rules for deciding the examples to keep from the majority class
 - e.g., Condensed Nearest Neighbors

Build the model (training) – induction

- **training set**
 - used to train the model

- **validation set** (extracted from the training data) can be used for optimizing the model
 - used to validate the model
 - on data unseen during training
 - unbiased evaluation of a model fit on the training dataset while tuning the model's hyperparameters.
 - **cross-validation** may be used
 - “best” model

Evaluate/test the performance of the model

- **test set**
 - unbiased evaluation of a final model fit on the training data set
- apply the model – *deduction*
- for evaluating the performance of the model, a cross-validation (CV) testing methodology should be used
- **CV**
 - helps in detecting overfitting
 - a way to test robustness
 - a statistical analysis of the results obtained during the CV (e.g., **confidence intervals - CIs**)
 - Central Limit Theorem for Cross-validation
 - given a sufficiently large sample size, the sampling distribution of the mean for a variable will approximate a normal distribution
 - perform multiple independent (or close to independent) evaluations of a model performance to result in a population of estimates
 - the mean of these estimates will be an estimate (with certain error) of the true underlying estimate of the model skill on the problem
 - around 30–50 repetitions for computing **CIs**
 - multiple (training-validation-testing) splits are used
 - **k-fold cross-validation**
 - k buckets
 - $k-1$ used for training, the remaining for testing
 - repeat k times
 - k performance measures P_1, P_2, \dots, P_k
 - μ - the mean performance
 - $$\mu = \frac{\sum_{i=1}^k P_i}{k}$$
 - compute the **CI** of μ
 - $[\mu - \alpha, \mu + \alpha]$ (or $\mu \pm \alpha$)
 - α is the *confidence value*
 - e.g., for 95% CI

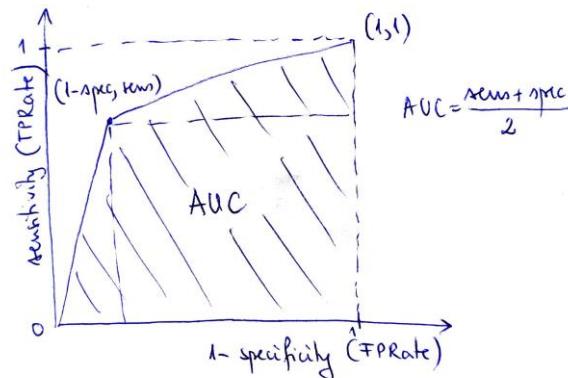
$$\alpha = 1.96 \cdot \frac{\sigma}{\sqrt{k}}$$

- 10 is usually used for k
- leave-one out cross-validation
 - $n-1$ instances for training, the remaining for testing
 - repeat n times
 - only one performance measure is provided
 - accuracy
 - computationally expensive
- “2/3-1/3” split
 - randomly select 2/3 from the data for training, the remaining for testing
 - repeat k times
 - CI of the mean performance
- [performance measures](#) used for assessing the performance of a ML model:
 - for classification:
 - accuracy, precision, recall (sensitivity), f-measure (f-score), specificity, AUC (Area under the ROC curve), AUPRC (Area under the precision/recall curve)
 - [0,1]
 - binary classification
 - confusion matrix (computed on a testing data set)
 - TP - # of true positives
 - TN - # of true negatives
 - FP - # of false positives
 - FN - # of false negatives

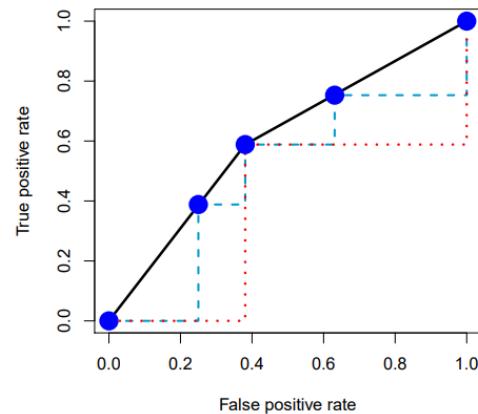
		True Class	
		Positive	Negative
Predicted Class	Positive	TP	FP
	Negative	FN	TN

$acc = \frac{TP + TN}{TP + TN + FP + FN}$	$prec_+ = \frac{TP}{TP + FP}$ (Positive predictive value, PPV)	$sens = recall_+ = \frac{TP}{TP + FN}$ (TPRate, probability of detection - POD)
$f-score_+ = \frac{2}{\frac{1}{prec_+} + \frac{1}{recall_+}}$	$prec_- = \frac{TN}{TN + FN}$ (Negative predictive value, NPV)	$spec = recall_- = \frac{TN}{TN + FP}$ (TNRate)
$f-score_- = \frac{2}{\frac{1}{prec_-} + \frac{1}{recall_-}}$	$f-score = \frac{f-score_+ + f-score_-}{2}$	Weigthed f-score

- higher values are better
- accuracy (acc)
 - n is the dimensionality of the testing data
 - $acc = \frac{1 + 0 + \dots + 1}{n}$
 - 95% CI of acc (the mean of a population parameter)
 - $\alpha = 1.96 \cdot \sqrt{\frac{acc \cdot (1 - acc)}{n}}$
- AUC
 - imbalanced data
 - only one point on the ROC (Receiver Operating Characteristic) curve



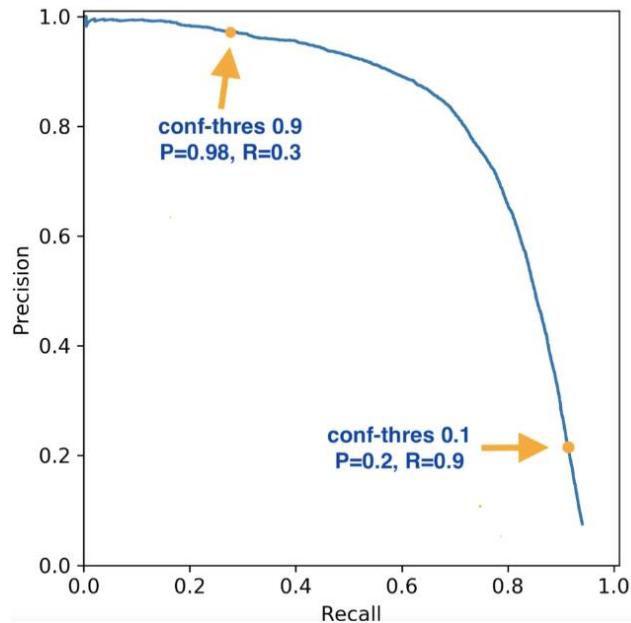
- multiple points



- performance of the classifier
 - AUC=0.5 random classifier
 - AUC $\in (0.5, 0.6]$ poor
 - AUC $\in (0.6, 0.7]$ fair
 - AUC $\in (0.7, 0.8]$ good
 - AUC $\in (0.8, 0.9]$ very good
 - AUC $\in (0.9, 1]$ excellent

- **AUPRC**

- Area under the Precision-Recall curve



- **other**

- e.g., [\(weather\) forecasting](#)
 - critical success index (CSI)
 - false alarm rate (FAR)
 - BIAS
 - HSS (Heidke Skill Score)
 - etc

- [multi-class classification](#)

- generalization from the binary case
- $M=(m_{i,j})_{i,j=1,n}$ where n is the number of classes
 - generalization from the binary case
 - $m_{i,j}$ = number of instances having the actual class j , predicted as belonging to class i
 - the evaluation measures are computed for each class i (e.g., $prec_i$, etc)
 - an aggregated measure is provided

- **for regression:** MAE (Mean of Absolute Errors), RMSE (Root Mean Squared Error), NRMSE (Normalized Root Mean Squared Error), R^2 (R-squared - coefficient of determination)
 - RMSE is more sensitive to outliers, whereas MAE treats all errors equally
 - MAPE (Mean absolute percentage error)
 - forecasting
 - MPE (Mean percentage error)

- lower values are better
 - e.g., NRMSE below 5%

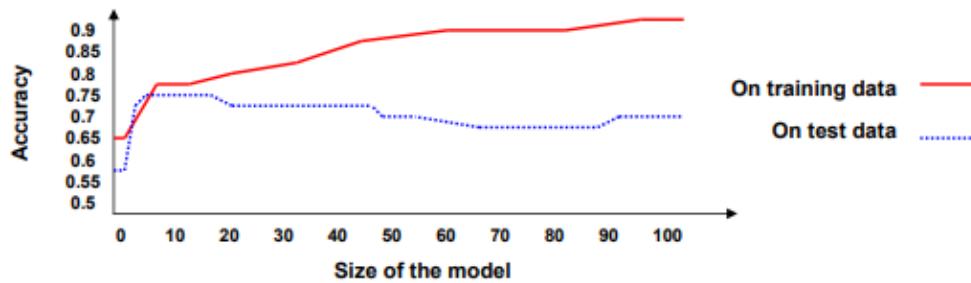
$MAE = \frac{\sum_{i=1}^n y_i - f(x_i) }{n}$ <p>y_i - predicted value for x_i $f(x_i)$ - true value (observation – in forecasting) n - # test instances</p>	$RMSE = \sqrt{\frac{\sum_{i=1}^n (y_i - f(x_i))^2}{n}}$
$MAPE = \frac{100\%}{n} \cdot \sum_{i=1}^n \left \frac{f(x_i) - x_i}{f(x_i)} \right $	$R^2 = 1 - \frac{SS_{RES}}{SS_{TOT}} = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2}$ <ul style="list-style-type: none"> • y – observations • \hat{y} – predictions

- comparing the performance of two classifiers/regressors
 - confidence intervals
 - statistical tests
 - is there a statistically significant difference between the performances of 2 models?
 - one/two-tailed [Wilcoxon signed rank test](#) for paired data, etc

Generalization

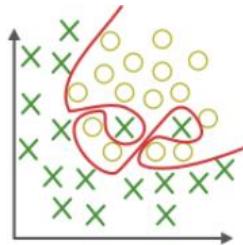
1. overfitting

- overfitting is finding complex functions due to
 - small number of training instances
 - noise or irrelevant data
- a hypothesis overfits the training data if a hypothesis that fits the training examples less well performs better on the entire data distribution
- an overfit model performs **well** on the **training set**, but usually has poor generalization capabilities



[2]

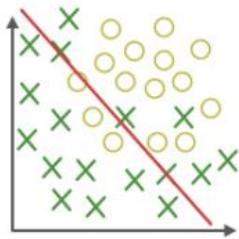
- causes
 - the model is too complex
 - models the training data too well
 - does not generalize
 - the training data is relatively small and is an insufficient representation of the underlying distribution that it is sampled from
 - the model fails to learn a generalizable mapping.



- good performance on training data but poor performance on testing data
- **techniques** to reduce overfitting:
 - Increase training data.
 - Reduce model complexity.
 - Early stopping during the training phase.
 - stop training once the model performance stops improving on a hold-out validation data set
 - Use regularization techniques (L2, L1 - [Ridge Regularization, Lasso Regularization](#))
 - reducing the generalization error without affecting the training error much
 - Use dropout for neural networks.

2. underfitting

- the model is too simple

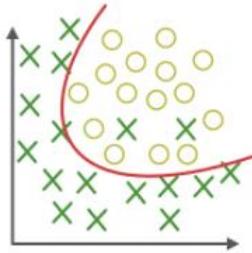


- a model that can neither model the training data, nor generalize to new data
- poor performance both on training data and on testing data

Overfitting and underfitting

- causes for poor performance of ML models

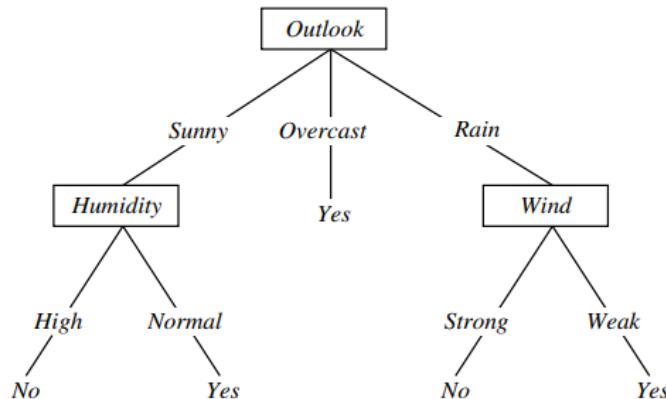
3. appropriate fitting



4. Decision tree (DT) learning

- **recent advances**
 - [Deep DTs \(Decision Streams\)](#)
 - deep directed acyclic graph of decision rules
 - [Deep neural DTs](#)
 - [Deep neural decision forests](#)
 - „Awesome DT research papers“ - <https://github.com/benedekrozemberczki/awesome-decision-tree-papers>
 - classification and regression tree research papers (before 2022)
 - prestigious conferences
 - *Machine Learning*: NeurIPS, ICML, ICLR
 - *Computer Vision*: CVPR, ICCV
 - *Natural Language Processing*: ACL, NAACL
 - *Artificial Intelligence*: AAAI, ICANN, IJCAI
 -
- used in **Data Mining** and **Machine Learning**, uses a decision tree as a predictive model
 - one of the practical methods for inductive inference
- it can be used both for **classification** and **regression**
- is an **eager** inductive learning model
- DT algorithms
 - ID3 – offline algorithm (Iterative Dichotomiser 3, Ross Quinlan in 1986)
 - ID4 is the online variant of ID3
 - C4.5 (C5.0) extensions of ID3
 - J48 is an open source Java implementation of C4.5 in Weka
 - CART (Classification And Regression Trees)
 - also
- follows the principle of **Occam's razor** in attempting to create the smallest DT possible
- **Limitations**
 - the problem of learning an optimal DT is NP-complete \Rightarrow the greedy algorithms lead to a local optimum
 - n boolean attributes/features, binary classification \Rightarrow how many distinct trees?
 - **overfitting** (complex trees)
 - the DT model is **unstable** (small fluctuations in data can make a large difference)
 - bagging, boosting

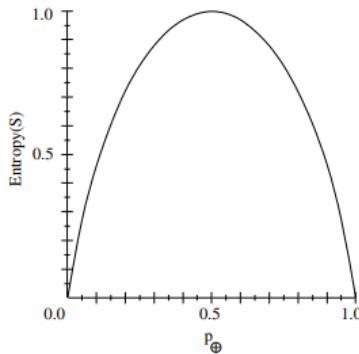
- **Advantages**
 - simple to understand and interpret
 - people are able to understand DT models
 - **interpretable/explainable ML (XML)**
 - making ML models and their decisions interpretable
 - uses a white box model
 - require little data preparation (other techniques often require data normalization, missing values to be removed)
 - able to handle both numerical and categorical data
 - **categorical variables**
 - fixed number of possible values
 - information can be sorted in categories
 - types of categorical variables
 - *ordinal* (some intrinsic order between the values, e.g. **rating**: excellent, good, fair, poor)
 - *nominal* – without any intrinsic order (e.g. **sex**: male, female)
 - *binary* (true, false)
 - **continuous variables**
 - can take an infinite number of possible values (e.g. real numbers)
 - e.g. temperature, blood pressure
 - perform well with large data sets
- **Decision tree representation**
 - each internal node tests an attribute
 - each branch corresponds to an attribute value
 - each leaf node assigns a classification
 - in DT learning, a hypothesis is a **decision tree**
 - DTs classify instances by sorting them down the tree from the root to some leaf node, which provides the classification of the instance



[1]

- DT – disjunctive expression
 - $(Outlook=Sunny \wedge Humidity=Normal) \vee (Outlook=Overcast) \vee (Outlook=Rain \wedge Wind=Weak)$

- **When to consider DTs**
 - instances describable by attribute-value pairs
 - target function is usually discrete valued (ID3/C4.5)
 - **disjunctive hypotheses** may be required
 - possibly noisy training data
 - training data may contain errors
 - missing attribute values in the training data
- **ID3 algorithm**
 - performs a **greedy** search through the space of possible decision trees \Rightarrow local optimum instead of the global one
 - recursive construction
 - select the “best” decision attribute for the current node
 - the attribute that is the most useful for classifying the examples
 - create descendants of the current node and build the subtrees recursively
 - for selecting the “best” attribute the **information gain** (IG) measure is used
 - measures the expected reduction in **impurity** due to sorting on the attribute
 - measures for **impurity**
 - **entropy**



- S is a sample of training examples
- p_{\oplus} is the proportion of positive examples in S
- p_{\ominus} is the proportion of negative examples in S
- Entropy measures the impurity of S

$$Entropy(S) \equiv -p_{\oplus} \log_2 p_{\oplus} - p_{\ominus} \log_2 p_{\ominus} \quad [1]$$

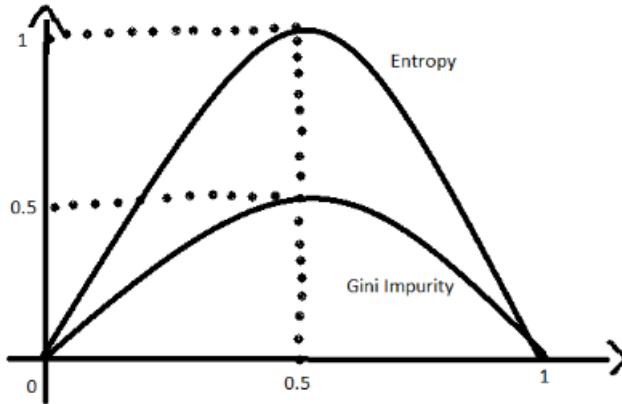
- c classes

$$E(S) = \sum_{i=1}^c -p_i \log_2 p_i$$

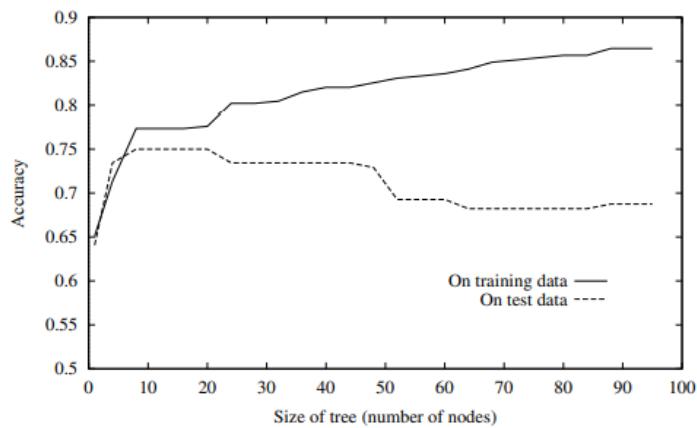
$$Gain(S, A) \equiv Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

- **Gini index**

$$Gini = 1 - \sum_{i=1}^n (p_i)^2$$



- about the same performance using entropy and Gini index
- **Inductive bias in ID3**
 - preference for short trees, and for those with **high information** gain near the root
- **Overfitting in DT learning**
 - long hypotheses (trees) may be obtained by coincidence



[1]

- avoiding overfitting by **pruning**
 - stop growing the tree when data split is not statistically significant
 - use χ^2 test to decide, when building the DT, if pruning should be applied at a certain node

- if the feature (attribute) at a certain node is uncorrelated with the decision of splitting the node, then prune
 - grow full tree, then post prune [1]
 - *reduced error pruning*
 - use a validation set
 - evaluate the impact on the validation set of pruning each possible node (plus those below it)
 - *rule post-pruning*
 - convert tree into rules
 - prune each rule independent of others
 - more successful in practice
 - used by C4.5
- **Issues in DT learning [1]** – handled by C4.5 algorithm
 - incorporating continuous valued attributes
 - alternative measures for selecting attributes
 - handling training examples with missing attribute values
 - handling attributes with costs

5. DT related research topics

- ensemble learning
 - Boosted Decision Trees
 - Random Forests (bagging)
- Fuzzy Decision Trees
- Lazy Decision Trees
- Hybrid models
 - DT + ANN (Artificial Neural Networks)
 - DT + SVM (Support Vector Machines)
 - DT + HMM (Hidden Markov Models)

[SLIDES]

- [Inductive learning and decision tree learning](#) (F. Leon) [2]
- [DTs](#) (T. Mitchell) [1]

[READING]

- [Decision tree learning](#) (T. Mitchell) [1]
- [Decision trees](#) (N. Nilsson) [2]

Bibliography

[1] Mitchell, T., *Machine Learning*, McGraw Hill, 1997 (available at www.cs.ubbcluj.ro/~gabis/ml/ml-books)

[2] Nilsson, N., *Introduction to Machine Learning*, Stanford University, 1996 (available at www.cs.ubbcluj.ro/~gabis/ml/ml-books)