

# Artificial neural networks (ANNs)

## SUMMARY

1. Artificial neural networks (ANNs) .....	1
2. Types of NNs for supervised learning .....	2
3. Deep Neural Networks (DNNs) .....	6
4. Convolutional Neural Networks (CNNs).....	7
5. Classical NN architectures .....	8
6. Other ANNs related research topics.....	20

## 1. Artificial neural networks (ANNs)


- Biological inspirations
  - Properties of the brain
    - it can **learn**, reorganize itself from experience
    - it **adapts** to changing conditions
    - it is **robust** and fault tolerant
- **Robustness** in ML?
  - the degree to which a model's performance changes when using new data
    - noise
  - ideally, performance should not deviate significantly
  - how to test it?
    - CV, CIs
- **Two types of learning in NNs**
  - **supervised**
    - FFNNs, RBFNs, RNNs, DNNs, CNNs, ...
  - **unsupervised**
    - self-organizing maps (SOM)
    - Hebbian learning
    - **autoencoders** (AEs)
      - **self-supervised learning**
      - **encoder-decoder** architecture
- **Characteristics** of supervised NN learning models
  - represent (complex) non-linear functions
  - **eager** inductive learning models
  - appropriate for **offline** and **online** learning
  - used for **classification** and **regression**

- **black-box** models
  - achieving **interpretability** – explainable neural models
- are **robust** to noisy data
- need for multiple and representative examples
- enable to model complex **static** phenomena (Feed-forward neural networks – FFNNs) as well as **dynamic** ones (e.g., Recurrent neural networks – RNNs)
  - **static** phenomena
    - time has no role
  - **dynamic** phenomena
    - temporal events
    - image and video recognition, time series, handwritten recognition, motion detection, signal processing, stock market prediction, speech recognition, also
- **NNs require**
  - a good representation of the data
  - training vectors must be statistically representative of the entire input space
  - the use of NNs needs a good comprehension of the problem
  - NNs require good *data preprocessing* (e.g., data normalization for numerical data)
    - the range of all features should be normalized
      - comparable range for the features
      - transpose the input variables into the range of the activation function codomain (i.e. for *logistic* [0, 1], for *tanh* [-1, 1])
    - speeds up learning, faster convergence
- Research domain: **NAS** (*Neural Architecture Search*)
  - subfield of **automated machine learning** (**AutoML**)
    - process of automating the tasks of applying ML to real-world problems
  - technique for automating the design of ANNs (both classical and deep)
    - **search space** defines the type(s) of ANN that can be designed and optimized
    - **search strategy** defines the approach used to explore the search space.
    - **performance estimation strategy** evaluates the performance of a possible ANN from its design (without constructing and training it).
  - RL, Hill climbing, Evolutionary algorithms, PSO, Multi-objective optimization,...

## 2. Types of NNs for supervised learning

### **1. Feed-forward neural networks (FFNNs)**

- an ANN where connections between the units do not form a directed cycle
- the first and simplest form of ANN
- the information moves in only one direction, forward, from the input nodes through the hidden nodes and the output nodes
- there are **no cycles** or **loops** in the network (time has no role)




(a) Feedforward network

- types of FFNNs

- **Multilayer perceptron (MLP)**


- consists of multiple layers of nodes in a directed graph, with each layer fully connected to the next one



- except for the input nodes, each node is a *neuron* (processing element) with a *nonlinear* activation function
- utilizes a supervised learning technique called *backpropagation* for training the network
- MLP is a modification of the standard linear perceptron and can distinguish data that are not linearly separable

- **Time delay neural networks (TDNNs)**


- [theory](#)
- an alternative to an NN architecture whose purpose is to work on continuous data
- learning a **temporal** sequence of events
- maps a finite time sequence  $\{X(t), X(t - \Delta), X(t - 2 \cdot \Delta) \dots X(t - m \cdot \Delta)\}$  into a single output  $y$  (this can be generalized for the case when  $x$  and/or  $y$  are vectors)



- [Pytorch](#)
- helpful in many applications like:
  - time series predictions
  - online spell check
  - speech recognition (generation)
  - image analysis
  - aso
- [Deep TDNN](#)

- **Radial basis function networks (RBFNs)**

- specific feed-forward architecture
- 1 hidden layer
- Gaussian activation function at the hidden layer




- connected to the **Instance based learning** (IBL) literature, but **eager** instead of **lazy**
  - computes a global approximation to the target function  $f$ , in terms of a linear combination of local approximations (“**kernel
  - is a different kind of two-layer neural network
    - i. the hidden units compute the values of kernel functions (local approximations)
    - ii. the output unit computes  $f$  as a linear combination of kernel functions**

- applications
  - [fault diagnosis](#)
  - forecasting
  - [image classification](#)
  - [image reconstruction](#)
  - aso

## 2. Recurrent neural networks (RNNs)

- **sequential** or time series data
  - RNNs are a variant of the conventional FFNNs that can deal with sequential data and can be trained to hold knowledge about the past.
    - a mechanism is required to retain past or historical information to forecast future values.
  - connections between units form a **directed cycle**
    - this creates an internal state of the network which allows to exhibit **dynamic temporal behavior**
    - can model systems with internal state (dynamic ones)



(b) Recurrent network [1]

- unlike FFNNs, RNNs can use their **internal memory** to process arbitrary sequences of inputs
- appropriate for time series data
  - learning is **sequential**
- applications:
  - handwritten recognition
  - motion detection
  - signal processing
  - text generation
  - time series prediction
  - stock market forecasting
  - aso
- the vanishing gradient problem of RNNs cause the network not to learn much → specialised versions of RNN
  - **LSTM**
  - **GRU (Gated Recurrent Unit)**

### Long-Short Term Memory networks (LSTMs)

- a type of RNN
- this model is an attempt to allow the unit activations to retain important information over a much longer period of time
- handling the *vanishing* and *exploding* gradient problem
  - includes specialized "gates" within its structure to manage long-term and short-term memory
- **applications:**
  - language learning
  - robot control
  - music composition
  - speech and handwriting recognition
  - video processing
  - ...
- other architectures: DeepLSTM, ConvLSTM, BiLSTM (Bidirectional LSTM), ensemble of LSTMs

### 3. Deep Neural Networks (DNNs)

- multiple hidden layers
- can express easier complex functions
- a layer may be viewed as a "*feature hierarchy*"
- Classes of DNNs [4]
  - DNNs for supervised learning
  - DNNs for unsupervised or generative learning
    - generative models – can learn and mimic any distribution of data
      - Boltzmann Machines, Restricted Boltzmann Machines, Deep Belief Networks, Deep Boltzmann Machines
        - Diffusion models
          - probabilistic models
          - parameterized Markov chain trained using variational inference (VI) to produce samples matching the data after finite time
          - VI – approximating of probability densities through optimization
          - Image Synthesis, Video Generation, 3D Generation, Medical Analysis, Text Generation, Speech Synthesis, Time Series Generation, Molecule Design, Graph Generation, etc
        - Diffusion classifiers
    - *Generative adversarial networks (GANs) [3]*
      - two nets competing one against the other (generator/discriminator)
        - learn to generate new data
          - generating images, faces, photographies
      - bidirectional GAN
      - various types of GAN models

- Conditional GAN, Cycle GAN, Convolutional GAN, etc
  - generative adversarial exploration for [reinforcement learning](#)
- [Variational AEs](#) (VAEs)
  - probabilistic latent space
  - provides a statistical manner for describing the samples of the dataset in latent space
  - latent code generated by the encoder is a probabilistic encoding
    - VAEs express not just a single point in the latent space but a distribution of potential representations
- Generative models in [reinforcement learning](#)
- **research:** solving unsupervised learning problems with DNNs (e.g. ICA – independent component analysis, feature analysis, aso)
- Hybrid DNNs, ensemble of DNNs, fuzzy DNNs
- [Graph Neural Networks](#) (GNNs)
  - a type of DL model that can be used to learn from graph data



## 4. Convolutional Neural Networks (CNNs)

- inspired by the organization of the visual cortex (biological inspiration)
- applications:
  - [computer vision](#) [3]
  - natural language processing
    - e.g., [sentence classification](#)
  - video processing
  - object detection and recognition
- are deep FFNNs
  - **convolution**
    - from a dimension of an input, a filter is applied to it to take some of the interesting features from that dimension
- **GCN** - graph convolutional networks
  - graph-structured data
  - handle higher dimensional (non-grid) data
  - **applications**
    - [semi-supervised learning](#)
    - supervised learning ([text classification](#))
    - [unsupervised learning](#)
- **MobileNets** – efficient CNN architecture for mobile devices
- **low resource devices**
  - **distillation**
    - compressing the knowledge from a large network into a smaller one
    - distilling the knowledge in a [NN](#) (Hinton, 2015), distilling knowledge from [GCN](#)
- Ensemble of CNNs

## 5. Classical NN architectures

- **Artificial neuron**

- non-linear, parameterized function with restricted output range




- the output of the neuron is obtained by applying an (non-linear) **activation function**  $f$  on the linear combination of the neuron inputs

- **Activation functions**

- *Signum* – output range: -1, +1
  - does not have a derivative, undifferentiable in 0
  - perceptron

Signum


$$f(x) = \begin{cases} +1 & x > 0 \\ -1 & \text{otherwise} \end{cases}$$



- *Identity* – output range: (-∞, +∞)

Identity


$$f(x) = x$$



- *Sigmoid (logistic)* – output range: (0, +1)
  - e.g. predict a probability

Sigmoid


$$f(x) = \frac{1}{1 + e^{-x}}$$



- *Hyperbolic* – output range: (-1, +1)
  - smooth approximation for the perceptron function
  - learning smoother than the perceptron

Hyperbolic

$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



- *Gaussian* – output range: (0, 1)
  - Radial basis function networks (RBFN)

Gaussian

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} \cdot e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



- *ReLU* (Rectified Linear Unit) – output range: (0,  $\infty$ )
  - undifferentiable in 0
  - main advantages
    - it does not activate all the neurons at the same time
    - converge very fast
  - problem
    - “dying”/“dead” ReLU
    - neurons becoming inactive during training and consistently outputting zero  $\Rightarrow$  loss of learning capability
  - deep architectures

ReLU

$$f(x) = \max(0, x)$$




- *Leaky ReLU*
  - introducing a small slope for negative values in the activation function  $\Rightarrow$  small gradient for negative inputs, preventing complete inactivity
  - allow a small, non-zero gradient when the unit is not active

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{otherwise} \end{cases}$$


- *PReLU* (Parametric Rectified Linear Unit) – output range: (- $\infty$ ,  $\infty$ )
  - improves ReLU and Leaky ReLU
    - **idea:** making the coefficient of leakage (0.01) into a parameter that is learned along with the other neural network parameters:

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{otherwise} \end{cases}$$

- *Soft plus* – output range:  $(0, \infty)$




- *eLU* (Exponential Linear Unit)
- *GeLU* (Gaussian Error Linear Unit)




### • Perceptron


- represents a hyperplane decision surface in the high dimensional space of instances



- binary classification (outputs: -1, +1)



- (a) **linearly separable** data set (i.e., data set can be separated by a straight line)
- (b) the classes are **not** linearly separable
- o linear classifier
- o appropriate for online learning



[1]

$$o(x_1, \dots, x_n) = \begin{cases} 1 & \text{if } w_0 + w_1 x_1 + \dots + w_n x_n > 0 \\ -1 & \text{otherwise.} \end{cases}$$

### Learning task

- **training examples**  $D = \{(x_d, t_d)\}_{d=1,s}$   $x_d = (x_{1d}, x_{2d}, \dots, x_{nd}) \in \Re^n$ ,  $t_d \in \{-1, +1\}$
- **goal**
  - learn the separating hyperplane
  - hypothesis:  $w = (w_0, w_1, \dots, w_n) \in \Re^{n+1}$
- **error function**
  - **online learning**  
 $E_d(\vec{w}) = t_d - o_d$
  - **offline learning**  
 $E(\vec{w}) = \frac{\sum_{d=1}^s |t_d - o_d|}{s}$

- o weights initialization
  - small random values (or 0)

- o **Training rule**

$$w_i \leftarrow w_i + \Delta w_i$$

where

$$\Delta w_i = \eta(t - o)x_i$$

Where:


- $t = c(\vec{x})$  is target value
- $o$  is perceptron output
- $\eta$  is small constant (e.g., .1) called *learning rate* [1]

- **Linear classifier**

$$o(\vec{x}) = \begin{cases} 1 & \text{if } \vec{w} \cdot \vec{x} > 0 \\ -1 & \text{otherwise.} \end{cases}$$

- [Example](#) - perceptron


- The perceptron is able represent some useful boolean function: AND, OR,  $\neg$ AND,  $\neg$ OR
- converges only if training data is linearly separable and the learning rate is sufficiently small (e.g., 0.1)
  - [Perceptron convergence theorem](#) Rosenblatt
    - for a finite set of linearly separable labeled examples, after a finite number of iterations, the algorithm yields a vector  $w$  that classifies perfectly all the examples.
  - XOR function is not representable using a perceptron  $\Rightarrow$  we need multilayered networks



**XOR function**

- **Linear unit**

- consider a *linear unit*, whose output  $o$  is 
$$o = w_0 + w_1 x_1 + \cdots + w_n x_n$$



### Learning task

- **training examples**  $D = \{(x_d, t_d)\}$   $x_d = (x_{1d}, x_{2d}, \dots, x_{nd}) \in \Re^n$ ,  $t_d \in \Re$ 
  - $t_d$  represents the output of the neuron for the input instance  $d$
- **goal**
  - learn the weights that minimize the squared error
  - (1) over the entire training samples  $D$

$$E[\vec{w}] \equiv \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$

- e.g., using *gradient descent optimization* for minimizing the error => **batch mode gradient descent**

(2) for each training sample  $d \in D$

$$E_d(\vec{w}) = \frac{1}{2} (t_d - o_d)^2$$


- e.g., using *gradient descent optimization* for minimizing the error => **incremental (stochastic) gradient descent**
- online learning

- **Gradient descent**

- the gradient of a function is a vector of *first derivatives* taken with respect to its constituent variables

$$\nabla f(p) = \begin{bmatrix} \frac{\partial f}{\partial x_1}(p) \\ \vdots \\ \frac{\partial f}{\partial x_n}(p) \end{bmatrix}$$

- the gradient specifies the direction that produces the steepest increase in E
  - $-\nabla E[\vec{w}]$  - the direction of steepest descent
- e.g., the **gradient (slope)** of a line shows how steep it is



Gradient

$$\nabla E[\vec{w}] \equiv \left[ \frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$

Training rule:

$$\Delta \vec{w} = -\eta \nabla E[\vec{w}]$$

i.e.,

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

$$\begin{aligned} \frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \frac{1}{2} \sum_d (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_d \frac{\partial}{\partial w_i} (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_d 2(t_d - o_d) \frac{\partial}{\partial w_i} (t_d - o_d) \\ &= \sum_d (t_d - o_d) \frac{\partial}{\partial w_i} (t_d - \vec{w} \cdot \vec{x}_d) \\ \frac{\partial E}{\partial w_i} &= \sum_d (t_d - o_d) (-x_{i,d}) \end{aligned}$$

[1]

- Training rule for
  - **batch mode gradient descent**

$$\Delta w_i = \eta \cdot \sum_{d \in D} (t_d - o_d) \cdot x_{i,d}$$


- **stochastic gradient descent** (Delta rule)

$$\Delta w_i = \eta \cdot (t_d - o_d) \cdot x_{i,d}$$

- linear unit training rule using gradient descent is guaranteed to converge
  - to hypothesis  $H$  with minimum squared error
  - given a small learning rate
  - even when training data contains noise
  - even when training data not separable by  $H$

- **Multilayer network**

- also known as Multilayer perceptron ([MLP \[3\]](#))
- can express a rich variety of non-linear decision surfaces



[1]

- NN learning
  - **eager**
    - learning model – **network**
    - **hypothesis** – the vector of weights
- [Example](#) of two 2-layer perceptron network for representing the XOR function
- **Gradient descent** (GD) training rule over a multilayer network is the **backpropagation** training algorithm
  - **idea**
    - For each training example
      - propagate the input forward through the network
      - propagate the errors backward through the network
    - derive gradient descent rules for training
      - one network unit (e.g. *sigmoid*)

- multilayer network of units → **backpropagation [1]**
- $D$  – training data
- **batch mode GD**
  - The learning rule is applied after all training instances are provided
$$E(\vec{w}) \equiv \frac{1}{2} \sum_{d \in D} \sum_{k \in \text{outputs}} (t_{kd} - o_{kd})^2$$
- **stochastic GD**
  - The learning rule is applied incrementally, after each training instance
$$E_d(\vec{w}) \equiv \frac{1}{2} \sum_{k \in \text{outputs}} (t_k - o_k)^2$$
  - E.g., for each output neuron  $j$ , the weight  $w_{ji}$  (from input  $i$  to unit  $j$ ) is updated with
$$\Delta w_{ji} = -\eta \frac{\partial E_d}{\partial w_{ji}}$$
  - E.g., if sigmoid output units
$$\frac{\partial E_d}{\partial \text{net}_j} = -(t_j - o_j) o_j (1 - o_j)$$

=>

$$\Delta w_{ji} = -\eta \frac{\partial E_d}{\partial w_{ji}} = \eta (t_j - o_j) o_j (1 - o_j) x_{ji}$$
  - $x_{ji}$  – the  $i$ -th input to unit  $j$
- **Loss (error) functions**
  - **regression**
    - Mean Squared Error (MSE) – L2 loss
      - Sensitive to outliers
    - Mean of Absolute Errors (MAE) – L1 loss
      - L1 – more robust and generally not affected by outliers
  - **classification**
    - **cross-entropy** loss
    - measures the error between two probability distributions
      - e.g., for binary classification (two classes – 0,1)
      - **binary cross entropy** loss
        - the average cross-entropy across all data samples

$$L = -\frac{1}{N} \left[ \sum_{j=1}^N [t_j \log(p_j) + (1 - t_j) \log(1 - p_j)] \right]$$

for  $N$  data points where  $t_i$  is the truth value taking a value 0 or 1 and  $p_i$  is the Softmax probability for the  $i^{th}$  data point.

- **Kullback-Leibler (KL) Divergence**
  - for comparing two data distributions
  - measures the difference between two probability distributions
    - used in unsupervised learning tasks where the objective is to uncover structure in data by minimizing the divergence between the true and learned data distributions
  - ML/Data sciences
    - VAEs optimization
    - GANs
    - **data model drifting**
      - **model drift** refers to the degradation of ML model performance due to changes in **data** or in the **relationships** between input and output variables (the underlying distributions of the features have changed over time)
  -

- **characteristics of backpropagation**

- GD over the entire network weight vector
- training is slow (eager model)
  - using network after training is fast
- will find a **local** error minimum (the error surface may contain many different local minima)
  - practice: run multiple times
  - weights initialization
    - Xavier initialization (using a Gaussian distribution)
- **momentum**


$$\Delta w_{ji}(n) = \eta \delta_j x_{ji} + \alpha \Delta w_{ji}(n-1)$$

- speed up the convergence of the network
- avoid convergence to a local minimum

- **Optimization** algorithms [3] used in NN learning

- GD
  - **first order** optimization algorithms
    - gradient
  - **stochastic GD**
    - extensions
      - **Adagrad** (Adaptive gradient algorithm)

- adaptive learning rate in optimization algorithm to reduce the needs of extensive hyperparameter fine-tuning
- **Adam** (Adaptive moment estimation)
  - a learning rate is maintained for each network weight (parameter) and separately adapted
  - adaptive learning rates
  - *momentum* adaptation
  - deep learning
- **RMSProp**
  - **Hinton, 2012**
  - speed up mini-batch learning
  - divide the learning rate for a weight by a running average of the magnitudes of recent gradients for that weight
- **minibatch GD**
  - performs an update for every batch with  $n$  training examples
- **second order optimization algorithms**
  - use the second derivative (the **Hessian**)
  - Hessian computations expensive
    - [Lazy Hessians](#)
- **optimize** the NN architecture
  - number of hidden layers, number of hidden neurons per layer, learning rate, momentum, etc
  - **genetic algorithms**
- **Problems** with gradient-based learning methods and backpropagation (the weights receive an update proportional to the partial derivative of the error function with respect to the weight)
  - **vanishing gradient**
    - in some cases, the gradient will be vanishingly small  $\Rightarrow$  preventing the weights in changing their values
      - classical activation functions such as **sigmoid** or **hyperbolic tangent** have gradients in  $(0,1)$
    - **solutions** to prevent the vanishing gradient problem
      - use other **activations functions** (whose derivative has a larger domain): ReLU, eLU, PReLU
      - use **residual networks** (ResNet)
        - [residual connections](#)/learning
          - a residual connection provides another path for data to reach latter parts of the neural network by skipping some layers.



- use **batch normalization** layers, normalize the input
- **exploding gradient**
  - the gradient is too large
  - the model became **unstable** and unable to learn from the training data
  - solutions to prevent the exploding gradient problem
    - fewer layers in the network
    - *clipping*
      - thresholding the value of the gradient
      - before performing the GD, assign a clip value if the gradient exceeds a threshold
      - **PyTorch, Tensorflow**
    - weight regularization (L1, L2)
- **Overfitting** in ANNs
  - may be due to
    - too many neurons (complex networks)
    - insufficient training data
    - not appropriate network architecture
      - it is not close enough to the problem context
  - reducing overfitting
    - use a **validation** set during training
    - **weight decay**
      - decrease weights with a small factor during each iteration
    - **regularization** techniques (L1, L2)
      - penalize large weights
      - add to the error function a regularization term

$$Loss = Error(y, \hat{y})$$

Loss function with no regularisation

$$Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N |w_i|$$

Loss function with L1 regularisation

$$Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N w_i^2$$

Loss function with L2 regularisation

- the hyperparameter has a value between 0 (no penalty) and 1 (full penalty)
  - *high* penalty => the model will underestimate the weights and underfit the problem
  - *low* penalty => the model will be allowed to overfit the training data
- **dropout**
  - randomly drop up neurons (with their connections) during training
  - deep networks
- **Underfitting** in ANNs
  - the model is too simple, it cannot capture the essence of the data
  - insufficient training, simplicity of the model, insufficient neurons
- Expressive capabilities of classical ANNs

Boolean functions:

- Every boolean function can be represented by network with single hidden layer
- but might require exponential (in number of inputs) hidden units

Continuous functions:

- Every bounded continuous function can be approximated with arbitrarily small error, by network with one hidden layer [Cybenko 1989; Hornik et al. 1989]
- Any function can be approximated to arbitrary accuracy by a network with two hidden layers [Cybenko 1988].

[1]

- shallow neural network (1 or 2 hidden layers)

## 6. Other ANNs related research topics

- Boosted ANNs
  - using a boosting algorithm for improving the performance of ANNs
- Ensemble of ANNs (LSTMs, Deep LSTMs)
- Fuzzy ANNs, Fuzzy Deep Neural Networks
- Lazy ANNs
- Hybrid models
  - ANN + DT
  - ANN + SVM (Support Vector Machines)
  - ANN for function approximation in Reinforcement Learning (RL)
- Parallel/Distributed ANNs
- Deep Residual Networks (ResNets), [Progressive Neural Networks](#), [Attention mechanism](#) [3]
- **Transformers**
  - Language models
  - Multi-head Attention
  - *Sequential data*
    - handling long dependencies between input sequence elements and enable parallel processing
  - Applications
    - **NLP**
      - QA ([VQA](#)), text summarization, language translation, etc
    - **Computer vision**
      - Image processing (classification, segmentation, translation, etc)
      - Vision Transformers ([ViT](#))
    - **Multi-modal learning**
      - Visual captioning, Text/Image/Video/Speech Generation
    - Classification tasks
      - [Time series classification](#)
    - [Semi-supervised](#) approaches
    - [Contrastive](#) learning
      - an approach that focuses on extracting meaningful representations by contrasting positive and negative pairs of instances

- ....

## [SLIDES]

- [Artificial neural networks](#) (T. Mitchell) [1]

## [READING]

- [Artificial neural networks](#) (T. Mitchell) [1]
- [Modern practical Deep networks](#) (Goodfellow et al.) [2]
- [CNNs](#) and [CNN architectures](#) (Zhang et al.) [3]
- [RNNs](#) and [Moderns RNNs](#) (Zhang et al.) [3]

## Bibliography

[1] Mitchell, T., *Machine Learning*, McGraw Hill, 1997 (available at [www.cs.ubbcluj.ro/~gabis/ml-books](http://www.cs.ubbcluj.ro/~gabis/ml-books))

[2] Ian Goodfellow, Yoshua Bengio, Aaron Courville, *Deep Learning*, MIT Press, 2016 (online edition at <http://www.deeplearningbook.org/>)

[3] Aston Zhang, Zachary C. Lipton, Mu Li, and Alexander J. Smola, *Dive into Deep Learning*, 2020 (<http://d2l.ai/>)

[4] Li Deng and Dong Yu, *Deep Learning. Methods and Applications*, Foundations and Trends® in Signal Processing, Volume 7 Issues 3-4, 2014 (<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/DeepLearning-NowPublishing-Vol7-SIG-039.pdf>)