

Support vector machines (SVMs)

SUMMARY

1. Support Vector Machines.....	1
2. Types of SVMs	4
3. Some kernel functions for SVMs.....	7
4. Overfitting.....	9
5. LIBSVM.....	10
6. SVM related research topics	10

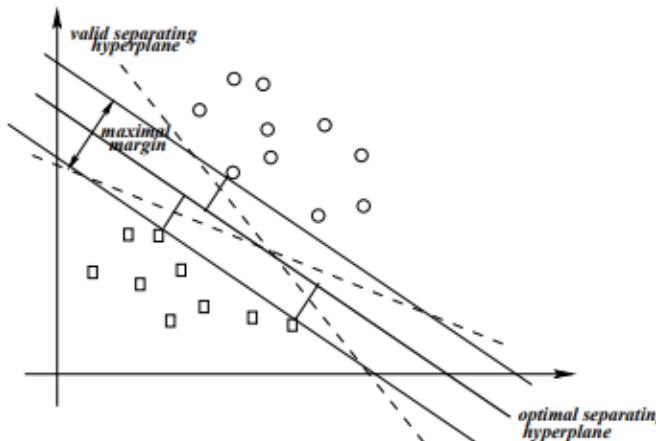
Performance of supervised learning models

- **predictive performance**
 - performance metrics (AUC,....,MAPE,)
- **computational complexity**
 - for the **training**, and **inference/testing** stages
 - time, space
 - *asymptotic analysis* (classes of complexity: O, θ,...)
 - *empirical analysis*
 - exact running time/required memory on specific data sets

1. Support Vector Machines

- SVMs are a set of related supervised learning methods used for classification and regression
- SVMs are eager inductive learners
- instances are high dimensional real valued vectors (data points in \Re^d)
- a SVM constructs a hyperplane or a set of hyperplanes in a high dimensional space, which can be used for classification, regression or other tasks.
 - a good separation is achieved by the hyperplane that has the largest distance to the nearest training data points of any class (the so called *functional margin*)
 - in general, the larger the margin, the lower the generalization error of the classifier \Rightarrow **large margin classifiers**

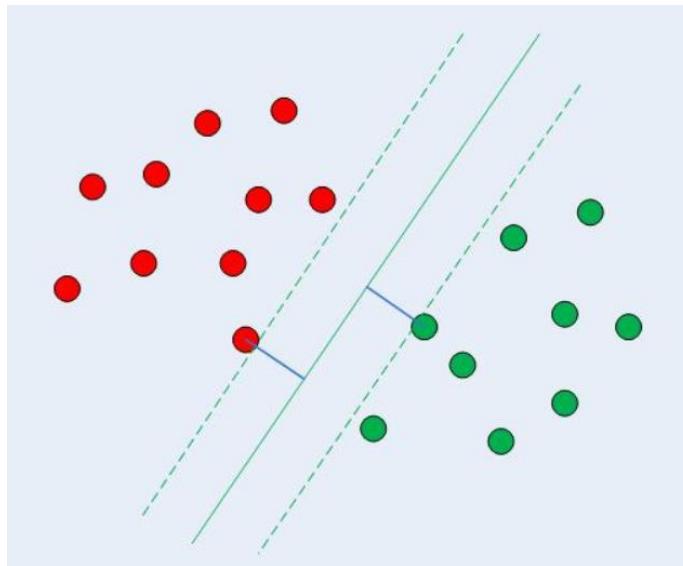
Optimal Separation Hyperplane



[6]

- SVMs – powerful ML models
 - [SVMs vs Pre-trained Language Models](#) (BERT, RoBERTa, DistilBERT) for Text Classification Tasks
- SVMs are inherently two-class classifiers
 - for multiple classes (M), two approaches have been proposed
 - **one-against-the-rest**
 - construct a hyperplane between class k and the $M-1$ other classes $\Rightarrow M$ SVMs
 - **one-against-one**
 - construct a hyperplane for any two classes $\Rightarrow M * (M-1)/2$ SVMs
- Research
 - **Deep SVMs**
 - Different perspectives
 - multiple layers of SVMs - [DSVM](#)
 - deep kernel learning architecture with multiple intermediate layers – [Totally Deep SVMs](#)
 - DSVMs for regression
 - [Transformers](#) as SVMs
 - a formal equivalence between the optimization geometry of self-attention and a hard-margin SVM problem that separates optimal input tokens from non-optimal tokens
 - [combining SVM and NN](#) by incorporating sample attention module
 - Applications
 - [SVMs with adaptive parameters](#) in financial time series forecasting
 - [Parallel SVMs](#)
 - [Max-margin contrastive learning](#)
 - Generative perspective
 - [VAE-based SVM detector](#)

- **Applications**
 - pattern recognition
 - classification (facial expression classification, image classification)
 - text categorization
 - e-mail filtering
 - web searching
 - sorting documents by topic
 - predictions (of traffic speed, protein structure prediction, breast cancer diagnosis prediction, time series prediction)
 - object detection, intrusion detection
 - handwritten recognition
 - ...
- **Main idea of SVMs**
 - Given a set of data points which belong to either of two classes
 - Find an optimal separating hyperplane
 - leaving the largest possible fraction of points of the same class on the same side
 - and
 - maximizing the distance of either class from the hyperplane



- Minimize the risk of misclassifying the training samples and the unseen test samples
- **Approach:** Formulate a *constrained optimisation problem*, then solve it using *constrained quadratic programming*
- **SMO** (Sequential Minimal Optimisation) algorithm, developed by John Platt, in 1986
- **Characteristics**
 - SVMs implement automatic complexity control to avoid overfitting
 - even if an SVM has a lot of hyperparameters, large margins make them simple classifiers

- intuition regarding the large margin:
 - points near the decision surface may represent very uncertain classification decisions
 - there is almost 50% chance of the classifier going either way
 - a classifier with large margin makes no very uncertain decisions
- SVR (Support Vector Regression)

2. Types of SVMs

- Linear SVMs
- Linear SVMs with soft margin
- Nonlinear SVMs

a). Linear SVMs

- the linear case
- data are linearly separable by a hyperplane
 - linear classifier
- formulate an optimisation problem
 - **Formalisation**
 - S a set of data points $x_i \in \Re^d$, $i=1,2,\dots,m$
 - two classes: +, -
 - the label of each instance x_i is $y_i \in \{-1,+1\}$
 - **training data** $\langle x_i, y_i \rangle$, $i=1,2,\dots,m$
 - impose a **functional margin** at least 1

The **set** S is **linear separable** if there are $w \in R^d$ and $w_0 \in R$ such that

$$y_i(w \cdot x_i + w_0) \geq 1 \quad i = 1, \dots, m$$

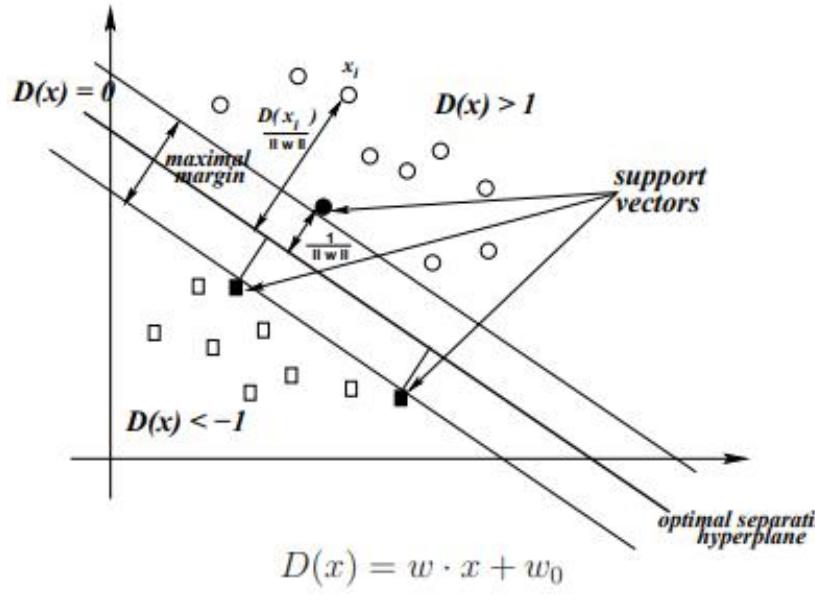
The pair (w, w_0) defines the hyperplane of equation $w \cdot x + w_0 = 0$, named the **separating hyperplane**.

[6]

The **s signed distance** d_i of a point x_i to the separating hyperplane (w, w_0) is given by $d_i = \frac{w \cdot x_i + w_0}{\|w\|}$.

It follows that $y_i d_i \geq \frac{1}{\|w\|}$, therefore $\frac{1}{\|w\|}$ is the **lower bound** on the distance between points x_i and the separating hyperplane (w, w_0) .

[6]



- **linear classifier**
 - $f(x) = \text{sgn}(w \cdot x + w_0)$
 - if $w \cdot x + w_0 > 0 \Rightarrow \text{predict 1, otherwise predict -1}$

- **Linear SVMs: the primal form**

$$\begin{aligned} &\text{minimize} \quad \frac{1}{2} \|w\|^2 \\ &\text{subject to} \quad y_i(w \cdot x_i + w_0) \geq 1 \quad \text{for } i = 1, \dots, m \end{aligned} \tag{1}$$
 - constrained quadratic problem (QP) with $d+1$ parameters
 - hypothesis \rightarrow hyperplane
 - $(w_0, w_1, \dots, w_d) \in \mathbb{R}^{d+1}$
 - if d is not very big (10^3) can be solved using **quadratic optimisation methods**
 - for large values of d (10^5)
 - Kuhn-Tucker theorem (nonlinear programming)
 - objective function and the associated constraints are convex
 - \Rightarrow the Lagrange multipliers ($\alpha_i \geq 0, i = 1, \dots, m$) are used to transform the problem into an equivalent **dual form**

- **Linear SVMs: the dual form**

$$\begin{aligned}
& \text{maximize} && \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j x_i \cdot x_j \\
& \text{subject to} && \sum_{i=1}^m y_i \alpha_i = 0 \\
& && \alpha_i \geq 0, i = 1, \dots, m
\end{aligned}$$

The link between the primal and the dual form:

The optimal solution (\bar{w}, \bar{w}_0) of the primal QP problem is given by

$$\begin{aligned}
\bar{w} &= \sum_{i=1}^m \bar{\alpha}_i y_i x_i \\
\bar{\alpha}_i (y_i (\bar{w} \cdot x_i + \bar{w}_0) - 1) &= 0 \text{ for any } i = 1, \dots, m
\end{aligned}$$

where $\bar{\alpha}_i$ are the optimal solutions of the above (dual form) optimisation problem.

[6]

b). Linear SVMs with soft margin

- the data set S is not linearly separable, or one ignores whether or not S is linearly separable
- extend the optimisation problem (1), by allowing a small number of misclassified points
 - better generalisation of computational efficiency
 - m non-negative variables ξ_i

- primal form

$$\begin{aligned}
& \text{minimize} && \frac{1}{2} \|w\|^2 + C \sum_{i=1}^m \xi_i \\
& \text{subject to} && y_i (w \cdot x_i + w_0) \geq 1 - \xi_i \text{ for } i = 1, \dots, m \\
& && \xi_i \geq 0 \text{ for } i = 1, \dots, m
\end{aligned}$$

- dual form

$$\begin{aligned}
& \text{maximize} && \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j x_i \cdot x_j \\
& \text{subject to} && \sum_{i=1}^m y_i \alpha_i = 0 \\
& && 0 \leq \alpha_i \leq C, i = 1, \dots, m
\end{aligned}$$

[6]

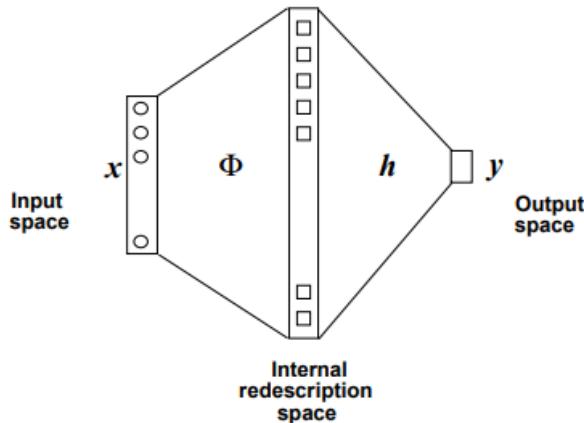
- introduce a parameter C – *misclassification parameter*
 - acts as a regularizing parameter
 - reducing overfitting
 - large $C \Rightarrow$ minimize the number of misclassified points
 - small $C \Rightarrow$ maximize the functional margin $\frac{1}{\|w\|}$

c). Nonlinear SVMs

- the data points from the input space \Re^d are mapped into a higher dimensional space \Re^n ($n > d$) using a function called a map $\phi: \Re^d \rightarrow \Re^n$
 - in a higher dimensional space, it is likely that a linear separator can be constructed
- the training algorithm would depend on the scalar (dot) product $\phi(x_i) \cdot \phi(x_j)$

- constructing (via ϕ) a separating hyperplane with maximum margin in the higher dimensional space yields a nonlinear decision boundary in the input space

General Schema for Nonlinear SVMs



[6]

- nonlinear SVMs make use of the “**kernel trick**”
 - the dot product is computationally expensive
 - a kernel function $K: \mathcal{R}^d \times \mathcal{R}^d \rightarrow \mathcal{R}$ such that $K(x_1, x_2) = \phi(x_1) \cdot \phi(x_2)$, where \cdot is the dot product
 - by using the kernel function, it is possible to compute the separating hyperplane without explicitly constructing the map ϕ

3. Some kernel functions for SVMs

- **Polynomial:** $K(x, x') = (x \cdot x' + c)^q$
 - **RBF (radial basis function):** $K(x, x') = e^{-\frac{\|x-x'\|^2}{2\sigma^2}}$
 - **Sigmoide:** $K(x, x') = \tanh(\alpha x \cdot x' - b)$
- $$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

Example Considering that the dimensionality of the input space is 2 ($d=2$), show that the polynomial function K previously defined, for $c=0$ and $q=2$ is a kernel function.

Proof. Assuming that $x=(x_1, x_2)$ and $y=(y_1, y_2)$, we have that

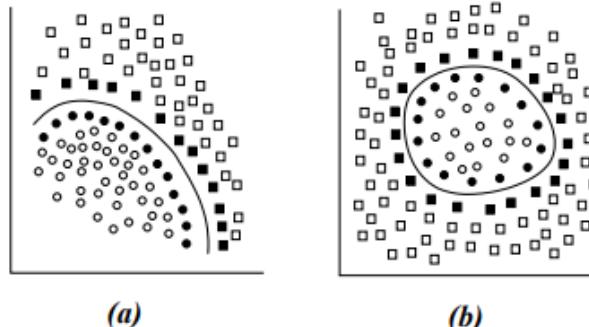
$$K(x, y)=(x \cdot y)^2=x_1^2 y_1^2+2 x_1 y_1 x_2 y_2+x_2^2 y_2^2$$

For proving that K is a kernel function, we must prove that there exists $\phi: \Re^2 \rightarrow \Re^n (n>2)$ such that $(x \cdot y)^2=\phi(x) \cdot \phi(y)$.

There are three possible mapping functions ϕ which satisfy the previous equality.

- a). $\phi: \Re^2 \rightarrow \Re^3, \quad \phi(x_1, x_2)=(x_1^2, \sqrt{2} x_1 x_2, x_2^2)$
- b). $\phi: \Re^2 \rightarrow \Re^3, \quad \phi(x_1, x_2)=1/\sqrt{2} (x_1^2-x_2^2, 2x_1 x_2, x_1^2+x_2^2)$
- c). $\phi: \Re^2 \rightarrow \Re^4, \quad \phi(x_1, x_2)=(x_1^2, x_1 x_2, x_1 x_2, x_2^2)$

- decision surfaces induced by a kernel function



Decision surface

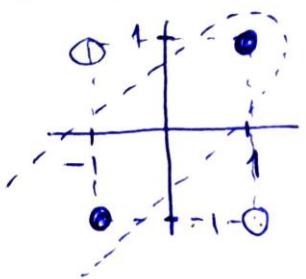
- (a) by a polynomial classifier, and
- (b) by a RBF.

Support vectors are indicated in dark fill.

[6]

XOR Example

SVM for the XOR problem



A	B	Output	
1	1	1	Class 1
-1	1	-1	Class 2
1	-1	-1	Class 2
-1	-1	1	Class 1

Polynomial Kernel

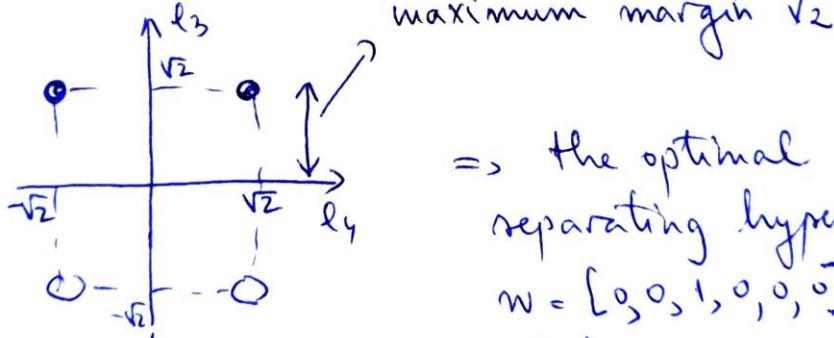
$$k(u, v) = (u \odot v + 1)^2 = (u_1 \cdot v_1 + u_2 \cdot v_2 + 1)^2 = u_1^2 v_1^2 + u_2^2 v_2^2 + 2u_1 u_2 v_1 v_2 + 2u_1 v_1 + 2u_2 v_2 + 1$$

$$\Phi(x_1, x_2) = (x_1^2, x_2^2, \sqrt{2}x_1 x_2, \sqrt{2}x_1, \sqrt{2}x_2, 1)$$

	x_1	x_2	l_1	l_2	l_3	l_4	l_5	l_6
Class 1	1	1	1	1	$\sqrt{2}$	$\sqrt{2}$	$\sqrt{2}$	1
	-1	-1	1	1	$\sqrt{2}$	$-\sqrt{2}$	$-\sqrt{2}$	1
Class 2	1	-1	1	1	$-\sqrt{2}$	$\sqrt{2}$	$-\sqrt{2}$	1
	-1	1	1	1	$-\sqrt{2}$	$-\sqrt{2}$	$\sqrt{2}$	1

=> points are separable on l_3 axis

maximum margin $\sqrt{2}$



=> the optimal
separating hyperplane

$$w = [0, 0, 1, 0, 0, 0] \quad (l_3=0)$$

and

$$w_0 = 0$$

- An SVM with a RBF kernel is equivalent with a 2 layered perceptron network

4. Overfitting

- C is a regularization parameter

- helps in avoiding overfitting
- theoretically, SVMs should be highly resistant to overfitting
 - in practice, it depends on the careful choice of C and other hyperparameters (e.g. the parameters of the kernel)
 - large $C \Rightarrow$ minimize the number of misclassified points
 - small $C \Rightarrow$ maximize the functional margin

5. LIBSVM

- implements the C-SVM algorithm
- for hyperparameters optimization (C and the parameters of the kernel) a grid search procedure is used
 - repeated trials for each parameter across a specified interval using geometric steps
 - for each combination of these parameters, a 10-fold cross-validation is performed during training, the quality of the combination being computed as the average of the accuracy rates estimated for each of the 10 divisions of the data
- the training step is computationally expensive

6. SVM related research topics

- [SVR](#) (regression)
- [Deep SVMs](#)
- [Deep SVMs for regression](#)
- Using SVMs in an ensemble learning
 - Boosted SVMs
 - Bagging SVMs
 - Stacking
 - base model
- Fuzzy SVMs
- Lazy SVMs
- Combining kernels for SVMs
- Hybrid models
 - [SVM + NN](#) (Artificial Neural Networks)
 - [SVM + kNN](#) (k -Nearest Neighbors)
 - [SVM + NBC](#) (Naive Bayes Classifier)

[SLIDES]

- [Introduction to Support Vector Machines](#) (Ming-Hsuan Yang and Antoine Cornuejols) [6]

[READING]

- [Support Vector Machines Explained](#) (T. Fletcher) [1]
- [Support Vector Machines](#) (Ch. 23) [2]
- [A Practical Guide to Support Vector Classification](#) (Hsu et al.) [3]

Bibliography

- [1] Tristan Fletcher, *Support Vector Machines Explained*, 2008, UCL
- [2] *Support Vector Machines*, Chapter 23, Cambridge University Press, 2012
- [3] Chih-Wei Hsu, Chih-Chung Chang, and Chih-Jen Lin, *A Practical Guide to Support Vector Classification*, Taiwan, 2009
- [4] [The Simplified SMO Algorithm](#)
- [5] John Platt, [Fast Training of Support Vector Machines using Sequential Minimal Optimization](#), Microsoft Research, USA