

# Instance based learning (IBL)

## SUMMARY


1. Instance based learning (IBL) .....	1
2. $k$ -Nearest Neighbor Learning ( $k$ NN).....	4
3. Distance-Weighted $k$ -Nearest Neighbor .....	6
4. Locally Weighted Regression (LWR).....	7
5. Radial Basis Function Networks (RBFN).....	8
6. Case Based Reasoning (CBR).....	9
7. IBL related research topics .....	9

## 1. Instance based learning (IBL)

- most of the methods from IBL are lazy methods
  - create local approximations of the target function
- **Main ideas** (lazy learning)
  - **training**: simply store all training examples
  - **at the query time**: compute only locally the target function
  - generalizing beyond the examples is postponed until a new instance must be classified
  - IBL methods are referred to “**lazy**” learning methods because the processing is delayed until a new instance is classified
    - instead of estimating the target function once for the entire instance space, the lazy methods can estimate it locally and differently for each new instance to be classified
- **Advantage**
  - IBL is useful in case of very complex target functions
- **Disadvantages**
  - can be computationally costly
    - the cost of classifying/evaluating an instance is high
    - almost all computations take place at the classification phase
  - usually considers all attributes
- **Main methods from IBL**
  - **$k$ -Nearest Neighbor ( $k$ NN)**
    - lazy learning
    - applications
      - computer vision (understanding and analysis of images), facial expression classification, object detection, text categorization, protein structure prediction, gene expression, etc
  - **Locally Weighted Regression (LWR)**
    - lazy learning
    - a generalization of kNN

- **Radial Basis Function networks (RBFN)**
    - combining IBL and neural networks
    - eager instead of lazy
  - **Case-based reasoning (CBR)**
    - symbolic representations and knowledge-based inference
    - lazy
- **kNN** and **RBFNs** are connected to
  - **local learning** - the main idea:
    - for each testing pattern
      - select the few training samples located in the vicinity of the testing pattern
      - train a classifier with only these few examples
      - apply the resulting classifier to the testing pattern
    - applications
      - e.g., [facial expression recognition](#)
  - **similarity-based learning**
    - based on computational measure of *similarity* (in the form of *distance* measure) between two instances from the input space  $X$
    - the *distance function*  $d:X \times X \rightarrow \mathbb{R}^+$ 
      - $d$  - **metric** function
        - non-negativity
        - coincidence axiom
        - symmetry
        - triangle inequality
      - $d$  - **semi-metric** function
        - non-negativity
        - coincidence axiom
        - symmetry
    - *similarity* – the inverse of the *distance*
    - example of distance functions
      - **metrics**: Euclidian, Minkovski, Manhattan, Levenshtein (for strings), Hamming, [Mahalanobis](#) (measures the distance between a multidimensional data point and a distribution)
      - **semi-metrics**: Cosine (used for texts), Pearson/Spearman (correlation)
    - **similarity learning**
      - learn the similarity between two samples from the input space
      - [Siamese Neural Networks](#)
        - type of *contrastive learning*
        - learn a similarity function
        - **applications**
          - signature verification
          - image recognition
          - [urban planning](#) (building change detection)
          - [recommendations](#)
          - [textual similarity](#)
          - [anomaly detection](#)

- a pair of neural networks that share weights and are designed to compute similarity functions, aiming to determine whether a pair of data is dissimilar or not
- mainly applied for images
  - [tabular data](#)



- we can use the learned function to differentiate between classes without needing a lot of data
  - [zero/one/few-shot learning](#) paradigm
    - building classifiers with little training data
  - [Siamese Transformer Networks](#) for Few-shot Image Classification
- useful for [representation learning](#)
  - extracting meaningful patterns from raw data to create representations that are easier to understand and process
  - these [representations](#) can be designed for interpretability, reveal hidden features, or be used for transfer learning
    - [supervised representation learning](#): learning representations on task A using annotated data and used to solve task B
    - [unsupervised representation learning](#): learning representations on a task in an unsupervised way (label-free data)
      - e.g., GPT, BERT

## 2. $k$ -Nearest Neighbor Learning ( $k$ NN)

- **When to consider  $k$ NN**
  - instances map to points in  $\Re^n$  ( $X = \Re^n$ )
  - a distance function  $d: X \times X \rightarrow \Re^+$  is defined
    - expresses the *distance* (dissimilarity) between two input instances
  - less than 20 attributes per instance
    - for large number of attributes, a dimensionality reduction should be applied (e.g. PCA, t-SNE)
  - lots of training data
  - the target function  $f$  to be learned is discrete or real-valued
- **Advantages**
  - training is fast
  - learn complex target functions
- **Disadvantages**
  - slow at query time
  - easily fooled by irrelevant attributes
- $k$ NN is a **lazy** learning method
  - **training:** simply store the training examples  $\langle x, f(x) \rangle$
  - **testing:** given a query instance  $y$ , we have to approximate  $f(y)$ 
    - **for classification**
      - $V = \{v_1, v_2, \dots, v_m\}$  – possible values for the target function
      - take a vote among the  $k$  nearest neighbors of  $y$  (from the training set)
      - the most common value of  $f$  among the  $k$  training examples nearest to  $y$

Given a query instance  $x_q$  to be classified,


- Let  $x_1, \dots, x_k$  denote the  $k$  instances from *training examples* that are nearest to  $x_q$
- Return

$$\hat{f}(x_q) \leftarrow \operatorname{argmax}_{v \in V} \sum_{i=1}^k \delta(v, f(x_i))$$

where  $\delta(a, b) = 1$  if  $a = b$  and where  $\delta(a, b) = 0$  otherwise.

[1]

- for instance, if binary classification and  $k=5$ , the class assigned to  $x_q$  is ‘-’



[1]

- **for regression**

- take the mean of the  $f$  values of the  $k$  nearest neighbors

$$\hat{f}(x_q) \leftarrow \frac{\sum_{i=1}^k f(x_i)}{k}$$

[1]

- **Characteristics**

- $k$ NN is simple, but works well in practice
- is among the top ten data mining algorithms
- the choice for  $k$  is **critical**
  - it depends on the data
  - $k=1 \Rightarrow$  nearest neighbor
  - a small value of  $k \Rightarrow$  noise will have a higher influence on the result
  - large  $k \Rightarrow$  reduces the noise, but is computationally expensive
  - if binary classification (i.e. two classes)  $\Rightarrow k$  must be an odd number
  - heuristics may be used for  $k$  – ex.  $k = \sqrt{n}$ , where  $n$  is the number of training samples
  - a good  $k$  can be obtained using **cross-validation**

- **Improving the results of  $k$ NN**

- preprocessing the training data (remove outliers, isolated points, normalized data)
- adapt metric to data
- learn the similarity function between instances  $\rightarrow$  **similarity learning**
- use **kernel methods**
  - methods that use only distances between objects, but no feature vectors
  - establish a distance measure between objects, then use only these distances

- **Optimization of  $k$ NN –  $kd$ -trees**

- the classification stage of  $k$ NN is slow
  - the retrieval of  $k$  nearest neighbors require  $n \log_2 n$ , where  $n$  is the number of training examples
    - assuming that  $k$  is a constant and a sorting algorithm in  $O(n \log_2 n)$  – e.g., MergeSort, HeapSort
- it is the eager variant of  $k$ NN

- *idea:* decrease the time to find the  $k$  nearest neighbors
- it has a slower training than  $k$ NN, but a faster classification
- train by constructing a lookup tree (***kd-tree***) – a balanced search tree, whose height is  $O(\log_2 n)$ 
  - each leaf nodes stores a training instance
  - nearby instances are stored at the same (or nearby) nodes
- the nearest neighbors will be found by searching the tree
  - require  $O(\log_2 n)$  steps
- **The curse of dimensionality problem**
  - when the data space is high dimensional, irrelevant attributes may dominate the  $k$ NN decision
  - **solution**
    - assign weights to the attributes
      - the relevant attributes have a larger weight
      - use a weighted distance, e.g. **weighted Euclidian distance**
    - use an approach similar to cross-validation to automatically choose values for the weights
    - a more drastic alternative is to eliminate the least relevant attributes from the instance space, i.e. setting the weight to 0.
- **$k$ NN behavior in the limit**
  - **$k$ NN** approaches the Bayes optimal learner, as the number of training instances  $\rightarrow \infty$  and  $k$  gets large
  - **1NN** approaches the Gibbs classifier, as the number of training instances  $\rightarrow \infty$

### 3. Distance-Weighted $k$ -Nearest Neighbor

- a variant of  $k$ NN in which the idea is to assign weights to the  $k$  nearest neighbors of the query instance
  - in classical  $k$ NN, all neighbors are weighted equally
- in DW- $k$ NN the nearer neighbors are weighted more heavily
  - **for classification**
    - the query instance is  $x_q$  and the  $k$ -nearest neighbors (from the training examples) of the query instance are denoted by  $x_1, x_2, \dots, x_k$

$$\hat{f}(x_q) \leftarrow \operatorname{argmax}_{v \in V} \sum_{i=1}^k w_i \delta(v, f(x_i))$$

$$w_i \equiv \frac{1}{d(x_q, x_i)^2}$$

[1]

- o **for regression**

$$\hat{f}(x_q) \leftarrow \frac{\sum_{i=1}^k w_i f(x_i)}{\sum_{i=1}^k w_i} \quad [1]$$

- $w_i = K(d(x_i, y))$ , where  $K(x) = \frac{1}{x^2}$
- $K$  is a kernel function (decreasing function)
- if  $k=n$  (we are using all training instances), the method is called **Shepard's method**
  - o it is a global method
  - o computationally costly

#### 4. Locally Weighted Regression (LWR)

- we note that  $k$ NN forms a local approximation to the target function  $f$  for each query point  $y$
  - **idea of LWR:** form an **explicit representation**  $\hat{f}(x)$  for the region surrounding the query point
    - o fit a function (e.g. **linear**, quadratic, multilayer neural net, etc) to  $k$  nearest neighbors
    - o produce a “piecewise approximation” to  $f$
    - o e.g. for locally weighted linear regression, if an instance is characterized by  $n$  attributes  $x = \langle a_1(x), a_2(x), \dots, a_n(x) \rangle$ , then
- $$\hat{f}(x) = w_0 + w_1 \cdot a_1(x) + w_2 \cdot a_2(x) + \dots + w_n \cdot a_n(x)$$
- o use a gradient descent approach for learning the weights, in order to minimize an **error function**
    1. squared error over  **$k$  nearest neighbors**
    2. distance weighted squared error over **all neighbors**
    3. distance weighted squared error over  **$k$  nearest neighbors**

If  $x_q$  is the query point:

1. Minimize the squared error over just the  $k$  nearest neighbors:

$$E_1(x_q) \equiv \frac{1}{2} \sum_{x \in k \text{ nearest nbrs of } x_q} (f(x) - \hat{f}(x))^2$$

2. Minimize the squared error over the entire set  $D$  of training examples, while weighting the error of each training example by some decreasing function  $K$  of its distance from  $x_q$ :

$$E_2(x_q) \equiv \frac{1}{2} \sum_{x \in D} (f(x) - \hat{f}(x))^2 K(d(x_q, x))$$

3. Combine 1 and 2:

$$E_3(x_q) \equiv \frac{1}{2} \sum_{x \in k \text{ nearest nbrs of } x_q} (f(x) - \hat{f}(x))^2 K(d(x_q, x))$$

[1]

- the gradient descent training rule for minimizing the error defined at 3 is

$$w_j \leftarrow w_j + \Delta w_j$$


where

$$\Delta w_j = \eta \sum_{x \in k \text{ nearest nbrs of } x_q} K(d(x_q, x)) (f(x) - \hat{f}(x)) a_j(x)$$

[1]

## 5. Radial Basis Function Networks (RBFN)

- compute a global approximation to the target function  $f$ , in terms of a linear combination of local approximations (“**kernel**” functions)
- is related to distance-weighted regression, but it is **eager**, instead of lazy
- is a different kind of two-layer neural network
  - the hidden units compute the values of kernel functions (local approximations)
  - the output unit computes  $f$  as a linear combination of kernel functions
- used for image classification, where the assumption of spatially local influences is well justified
- input instance  $x = \langle a_1(x), a_2(x), \dots, a_n(x) \rangle$



where  $a_i(x)$  are the attributes describing instance  $x$ , and

$$f(x) = w_0 + \sum_{u=1}^k w_u K_u(d(x_u, x))$$

One common choice for  $K_u(d(x_u, x))$  is

$$K_u(d(x_u, x)) = e^{-\frac{1}{2\sigma_u^2}d^2(x_u, x)}$$

[1]

- the two layers of the network are trained separately
- $k$  – a user-provided threshold specifying the number of kernel functions
- $x_u$  are prototype vectors/centers
  - can be selected using [clustering](#)


## 6. Case Based Reasoning (CBR)

- is instance-based learning applied to instance spaces  $X \neq \mathbb{R}^n$
- **characteristics** of [CBR](#) methods
  - are lazy learning methods
  - they classify new instances by analyzing similar instances while ignoring instances that are very different from the query
  - are similar to human reasoning
- instances are represented using a symbolic description
- a different distance metric is needed, adapted to the data
- **applications**
  - design of mechanical devices (CADET) [1]
  - scheduling problems
  - recommender systems
  - predator and prey problem
  - [QA](#)

## 7. IBL related research topics

- Fuzzy kNN

- Fuzzy CBR
- Fuzzy RBF
- Adaptive kNN
- Using kNN for improving AdaBoost
- Hybrid ML models
  - kNN+SVM (Support Vector Machines)
  - [Deep NN+RBFN](#) (for abnormality classification)
  - CBR + deep learning ([traffic management](#))
- **Deep RBFNs**
  - [Deep RBFN](#) for anomaly detection
  - GA based DRBFN for [medical classification](#)
  - Adaptive deep gradient RBFN for [industrial processes](#)
- **Deep kNN**
  - different formulations
  - [Hybrid model](#): kNN + DNN
  - DkNN for [medical diagnosis](#) classification
    - Feature extraction step (CNN)
    - A new loss function + neighbor search during training



- DkNN for [noisy labels](#)
- [KNN-enhanced Deep Learning Against Noisy Labels](#)
- [Deep Similarity-Enhanced kNN](#)
  - the similarity function is learned

## [SLIDES]

- [Instance based learning](#) (T. Mitchell) [1]

## [READING]

[Instance based learning](#) (T. Mitchell) [1]

## Bibliography

- [1] Mitchell, T., *Machine Learning*, McGraw Hill, 1997 (available at [www.cs.ubbcluj.ro/~gabis/ml/ml-books](http://www.cs.ubbcluj.ro/~gabis/ml/ml-books))