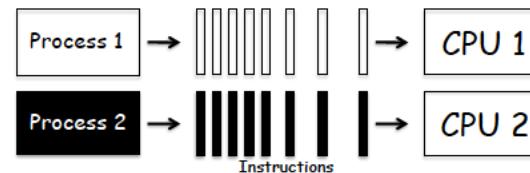


L2

Multiprocessing and Multithreading
C++ Threads

Multiprocessing, parallelism

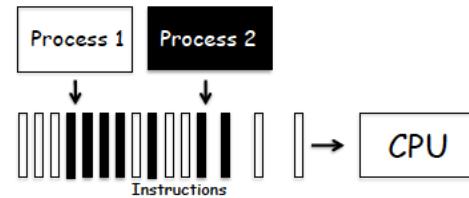
- Many of today's computations can take advantage of
- multiple processing units (through multi-core processors):



- Terminology:
- Multiprocessing : the use of more than one processing unit in a system
- Parallel execution: processes running at the same time

Multitasking, concurrency

- Even on systems with a single processing unit we may give the illusion of that several programs run at once
- The OS switches between executing different tasks



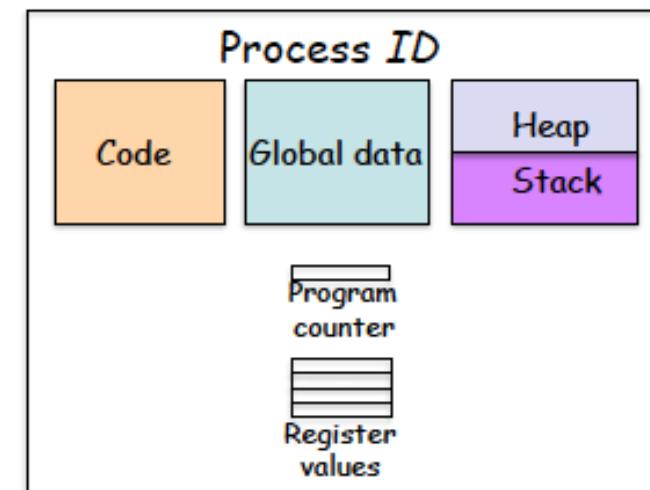
- Terminology:
 - Interleaving: several tasks active, only one running at a time
 - Multitasking: the OS runs interleaved executions
 - Concurrency: multiprocessing, multitasking, or any combination

Processes

- A (sequential) program is a set of instructions
- A process is an instance of a program that is being executed

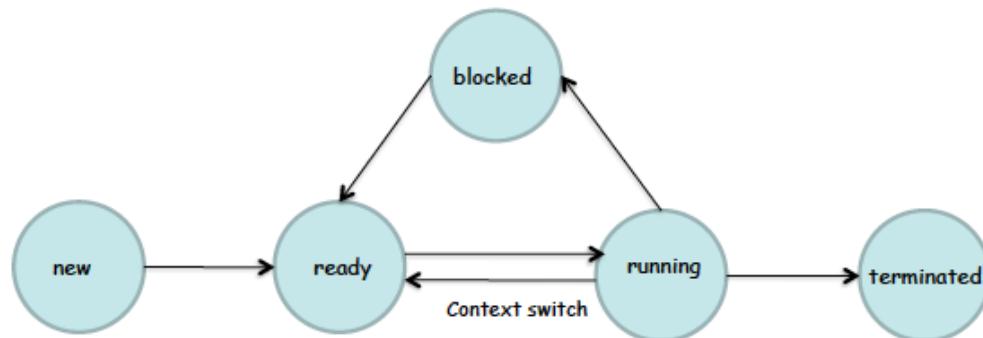
Operating system processes

- How are processes implemented in an operating system?
- Structure of a typical process:
 - Process identifier: unique ID of a process.
 - Process state: current activity of a process.
 - Process context: program counter, register values.
 - Memory: program text, global data, stack, and heap.



The scheduler

- A system program called the scheduler controls which processes are running; it sets the process states:
 - new: being created.
 - running: instructions are being executed.
 - blocked: currently waiting for an event.
 - ready: ready to be executed, but not been assigned a processor yet.
 - terminated: finished executing.

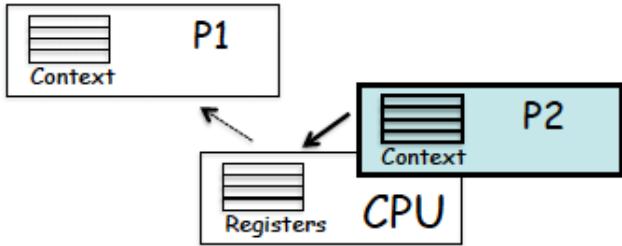


Blocked processes

- A process can get into state blocked by executing special program instructions
- When blocked, a process cannot be selected for execution
- A process gets unblocked by **external events** which set its state to ready again

The context switch

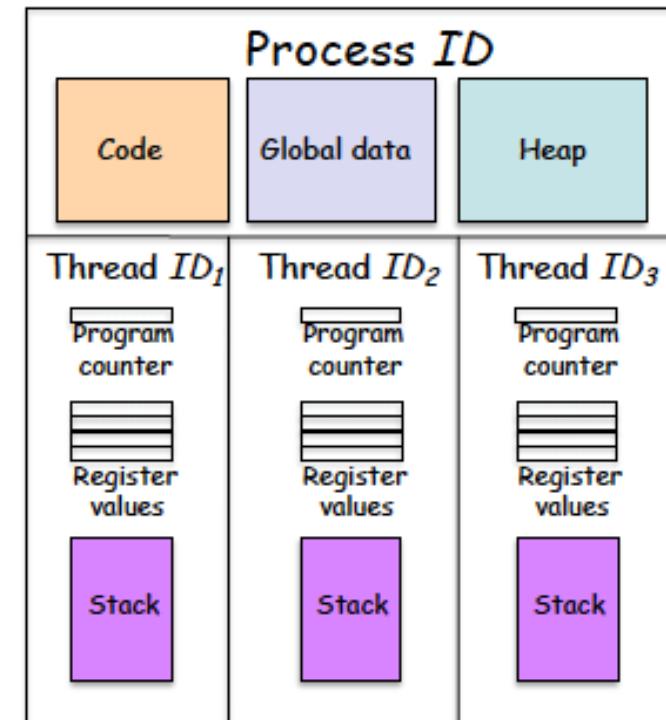
- The swapping of processes on a processing unit by the scheduler is called the context switch



- Scheduler actions when switching processes P1 and P2:
 - $P1.state := \text{ready}$
 - Save register values as P1's context in memory
 - Use context of P2 to set register values
 - $P2.state := \text{running}$

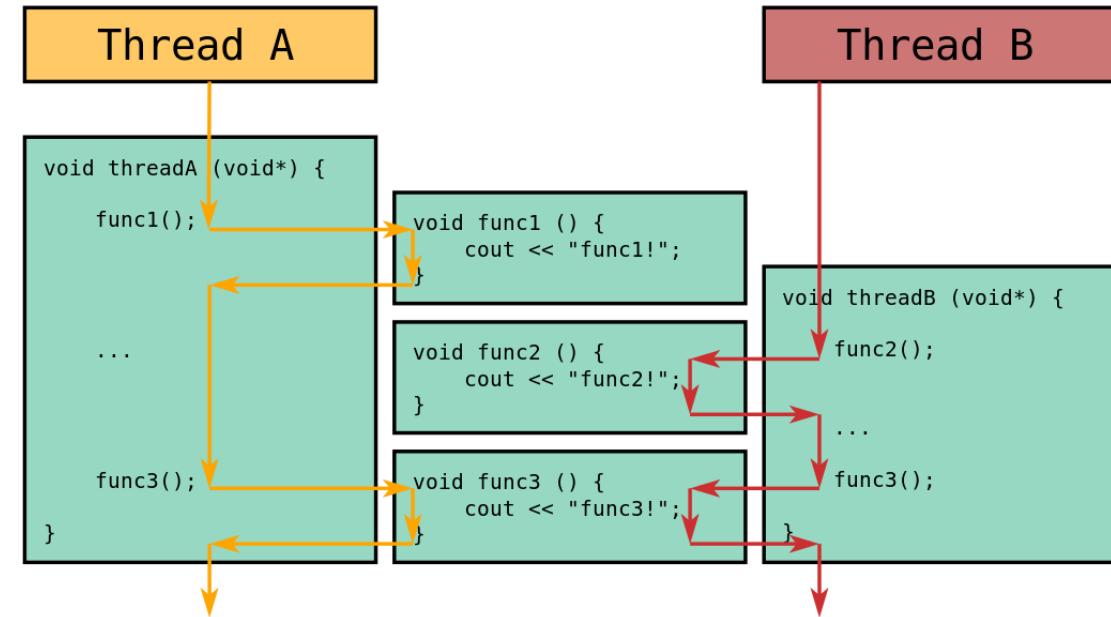
Threads

- Make programs concurrent by associating them with threads
- A thread is a part of an operating system process
- Components private to each thread
 - Thread identifier
 - Thread state
 - Thread context
 - Memory: only stack
- Components shared with other threads:
 - Program text
 - Global data
 - Heap



C++ threads (C++11 and further)

- `std::thread` class
- Each instance of this object
 - represents
 - wraps
 - managesa single execution thread.



Source: <https://kholdstare.github.io/technical/2012/08/21/objects-and-threads-in-cpp-1.html>

- all code is accessible to any thread
- two threads could be executing the same function at the same time.

Simple examples

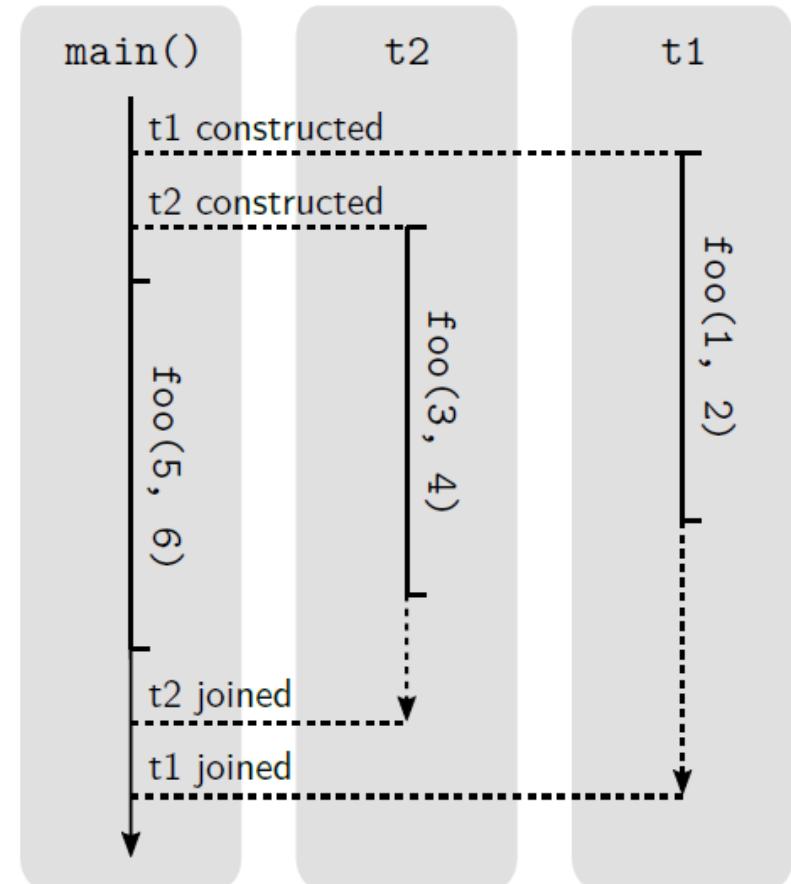
```
#include <iostream>
#include <thread>
void thread_function()
{
    std::cout << "thread function\n";
}
int main() {
    std::thread t(thread_function); // t starts running
    std::cout << "main thread\n";
    t.join(); // main thread waits for the thread t to finish
    return 0;
}
```

```
#include <iostream>
#include <thread>
#include <string>
void thread_function(std::string s)
{
    std::cout << "thread function ";
    std::cout << "message is = " << s << std::endl;
}
int main() {
    std::string s = "HPC&BDA";
    std::thread t(thread_function, s );
    std::cout << "main thread message = " << s << std::endl;
    t.join();
//t.detach();
    return 0;
}
```

Thread creation and destruction

```
#include <thread>  
  
void foo(int a, int b);  
  
int main(){  
    //Pass a function and args  
    std::thread t1(foo, 1, 2);  
    //Pass a lambda  
    std::thread t2([](){ foo(3,4); }); foo(5, 6);  
    t2.join();  
    t1.join();  
}
```

join is mandatory unless detach is called



Other useful functions

- `std::this_thread::sleep_for()`
: Stop the current thread for a given amount of time
- `std::this_thread::sleep_until()`
: Stop the current thread until a given point in time
- `std::this_thread::yield()`
: Let the operating system schedule another thread
- `std::this_thread::get_id()`
: Get the (operating-system-specific) id of the current thread
- `std::thread::detach()`
: Separates the thread of execution from the thread object, allowing execution to continue independently.

Sharing data

Global Variables

- All global and static variables that are initialized at compile time can be accessed by threads.
- Since the threads should know the addresses for them.

Passing By Value vs by Reference

- All parameters passed to a function when starting a thread are passed by value!
 - For passing by reference, we need to explicitly wrap the arguments in std::ref() .
 - Because the thread functions can't return anything, passing by reference is the only way to properly get data out of a thread without using global variables.
-
- Example:

```
void ref_function(int &a, int b) {  
    int val;  
    std::thread ref_function_thread(ref_function, std::ref(val), std::ref(val));  
}
```

static and thread_local Variables

```
void method() {  
    static int var = 0;  
  
    var++;  
}
```

- This does NOT create a separate instance of the static variable per thread instance !!!!
This is because static variables are initialized once when the compiler goes over their declaration.
- In order to have 'static' variables that are static within the scope of each particular thread, we must use `thread_local` variables instead.
Then each thread will have its own version of the static variable, and the static variable will only be destroyed on thread exit.

```
void method() {  
    thread_local int var = 0;  
  
    var++;  
}
```

race condition

- **race condition (race hazard)**
- *a condition of a program where its behavior depends on relative timing or interleaving of multiple threads or processes*
- The *race condition* term was introduced before 1955,
[e.g. David A. Huffman's doctoral thesis "The synthesis of sequential switching circuits", 1954]

Race condition

```
void threadfunc(unsigned* x) {  
//?????????????  
    for (int i = 0; i != 10000000; ++i) {  
        *x += 1;  
    }  
}
```

```
int main() {  
    std::thread th[4];  
    unsigned n = 0;  
    for (int i = 0; i != 4; ++i) {  
        th[i] = std::thread(threadfunc, &n);  
    }  
    for (int i = 0; i != 4; ++i) {  
        th[i].join();  
    }  
    printf("%u\n", n);  
}
```

A data race occurs when two or more threads access the same variable concurrently, and at least one of the accesses is a write.

critical race condition vs. non-critical race condition

- *critical race condition* = when the order may change the final state
- *non-critical race condition* = when the order do not change the final state

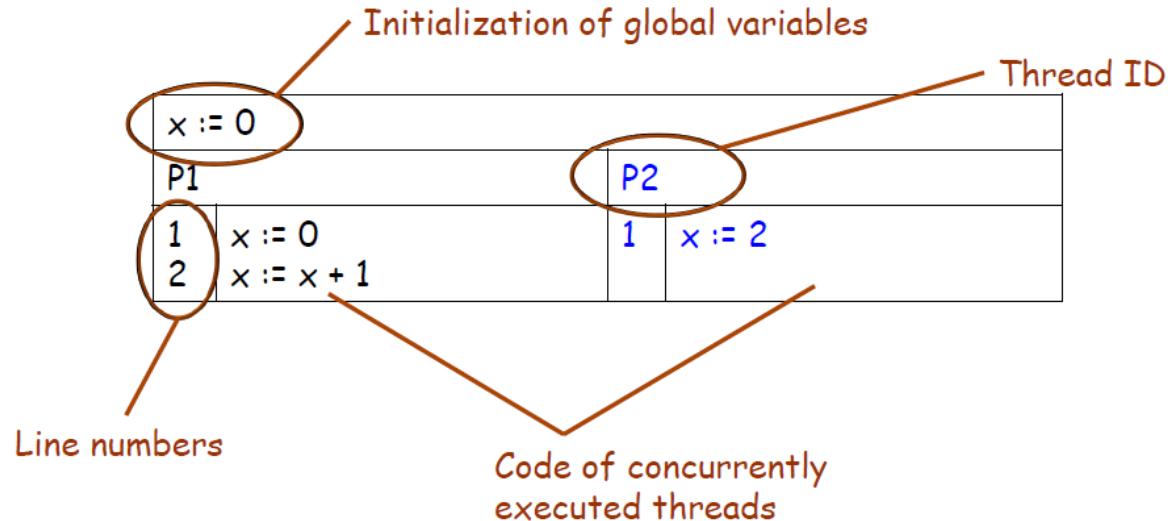
T1:

```
update (){  
    a=a+1;  
}
```

T2:

```
update (){  
    b=a*2;  
}
```

Concurrency



Execution Variants

P2	1	x := 2	x = 2
P1	1	x := 0	x = 0
P1	2	x := x + 1	x = 1

P1	1	x := 0	x = 0
P2	1	x := 2	x = 2
P1	2	x := x + 1	x = 3

P1	1	x := 0	x = 0
P1	2	x := x + 1	x = 1
P2	1	x := 2	x = 2

P1	1	x := 0	x = 0
P2	1	x := 2	x = 2
P1	2	x := x + 1	x = 3

Instruction executed with Thread ID and line number

Variable values after execution of the code on the line

Atomic instruction

- Levels of atomicity

Ex: $x := x + 1$

Execution:

temp := x

LOAD REG, x

temp := temp + 1

ADD REG, #1

x := temp

STORE REG, x

Variante de executie

x := 0			
P1		P2	
1	x := 0	1	x := 2
2	temp := x		
3	temp := temp + 1		
4	x := temp		

- "interleaving"

P1	1	x := 0	x = 0
P1	2	temp := x	x = 0, temp = 0
P2	1	x := 2	x = 2, temp = 0
P1	3	temp := temp + 1	x = 2, temp = 1
P1	4	x := temp	x = 1, temp = 1

Counter -> Details – register level

get this.count from memory into register

add value to register

write register to memory

- Example of interleaving

this.count = 0;

A: reads this.count into a register (0)

B: reads this.count into a register (0)

B: adds value 2 to register

B: writes register value (2) back to memory. this.count now equals 2

A: adds value 3 to register

A: writes register value (3) back to memory. this.count now equals 3

Race Conditions & Critical Sections

- A **Critical Section** = code segment where a data-race may occur

```
public class Counter {  
    protected long count = 0;  
    public void add(long value){  
        this.count = this.count + value;  
    }  
}
```

If an object Counter is used by
more threads!

=> Not *thread-safe*!

- Method `add()` is an example of critical section.

Atomics

- Processors have special, *atomic* instructions that always execute without racing with any other processor's accesses.
- These instructions are "indivisible, i.e., the processor does not internally decompose atomic instructions into smaller micro-code instructions.
- Atomic instructions are generally slower and more expensive (in terms of energy) to execute than normal instructions.

```
void threadfunc(std::atomic* x) {  
    for (int i = 0; i != 10000000; ++i) {  
        x->fetch_add(1);  
        // `*x += 1` and `(*x)++` also work!  
    }  
}
```

Uses processor lock-prefixed instructions !

Mutual exclusion - Mutexes

- **Mutual exclusion** means that at most one thread accesses the shared data at a time!

```
std::mutex mutex;
```

```
void threadfunc(unsigned* x) {  
    for (int i = 0; i != 10000000; ++i) {  
        mutex.lock();  
        *x += 1;  
        mutex.unlock();  
    }  
}
```



Where
should be
declared?

- A mutex (a kind of a data structure) has an internal state (denoted by state), which can be either locked or unlocked. The semantics of a mutex object is as follows:
- Upon initialization, state = unlocked.
- When a thread call:
 - `mutex::lock()` method: waits until state becomes unlocked, and then atomically sets state = locked. Note the two steps shall complete in one atomic operation.
 - `mutex::unlock()` method: asserts that state == locked, then sets state = unlocked.
- Binary semaphore

See also `std::unique_lock`
a RAII wrapper for exclusive locking

Recursive Mutexes

```
#include <mutex>

std::mutex mutex;

void bar() {
    std::unique_lock lock(mutex);
    // do some work...
}

void foo() {
    std::unique_lock lock(mutex);
    // do some work...
    bar(); // will deadlock
}
```

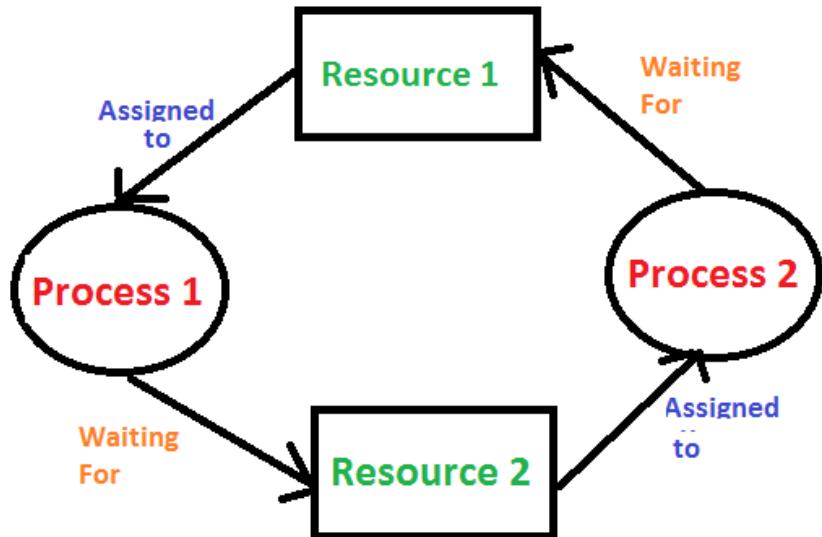
- The code will deadlock since `std::mutex` can be locked at most once

Recursive Mutexes

- `std::recursive_mutex` implements recursive ownership semantics
- The same thread can lock an `std::recursive_mutex` multiple times without blocking
- Other threads will still block if an `std::recursive_mutex` is currently locked
- Can be used with `std::unique_lock` just like a regular `std::mutex`
- Useful for functions that call each other and use the same mutex.

```
#include <mutex>  
  
std::recursive_mutex mutex;  
  
void bar() {  
    std::unique_lock lock(mutex);  
}  
  
void foo() {  
    std::unique_lock lock(mutex);  
    bar(); // OK, will not deadlock  
}
```

Deadlocks



```
std::mutex m1, m2, m3;  
void threadA() {  
// ???  
std::unique_lock l1{m1}, l2{m2}, l3{m3};  
}  
void threadB() {  
// ???  
std::unique_lock l3{m3}, l2{m2}, l1{m1};  
}
```

Possible deadlock scenario

1. `threadA()` acquires locks on `m1` and `m2`
2. `threadB()` acquires lock on `m3`
3. `threadA()` waits for `threadB()` to release `m3`
4. `threadB()` waits for `threadA()` to release `m2`

Solution

Deadlocks can be avoided by always locking mutexes in a *globally* consistent order

- Ensures that one thread always “wins”.
- Maintaining a globally consistent locking order requires considerable developer discipline.
- Maintaining a globally consistent locking order may not be possible at all.

```
std::mutex m1, m2, m3;  
  
void threadA() {  
    // OK, will not deadlock  
    std::unique_lock l1{m1}, l2{m2}, l3{m3};  
}  
  
void threadB() {  
    // OK, will not deadlock  
    std::unique_lock l1{m1}, l2{m2}, l3{m3};  
}
```

Condition variables

- A condition variable is a synchronization primitive that allows multiple threads to wait until an (arbitrary) condition becomes true.
- A condition variable uses a mutex to synchronize threads.
- Threads can *wait* on or *notify* the condition variable.
- When a thread waits on the condition variable, it blocks until another thread notifies it.
- If a thread waited on the condition variable and is notified, it holds the mutex.
- A notified thread must check the condition explicitly because *spurious wake-ups* can occur

`std::condition_variable` header <condition_variable> which has the following member functions:

- `wait()`: Takes a reference to a `std::unique_lock` that must be locked by the caller as an argument, unlocks the mutex and waits for the condition variable
- `notify_one()`: Notify a single waiting thread, mutex does not need to be held by the caller
- `notify_all()`: Notify all waiting threads, mutex does not need to be held by the caller

Worker threads example

```
#include <condition_variable>
#include <iostream>
#include <thread>

std::mutex a_mutex;
std::condition_variable condVar;
bool dataReady = false;

void waitingForWork(){
    std::cout << "Waiting\n";
    std::unique_lock<std::mutex> lck(a_mutex);
    condVar.wait(lck, []{ return dataReady; });
    std::cout << "Running\n";
}
```

```
void setDataReady(){
{
    std::lock_guard<std::mutex> lck(a_mutex);
    dataReady = true;
}
std::cout << "Data prepared\n";
condVar.notify_one(); }
```

```
int main(){
    std::thread t1(waitingForWork);

    std::thread t2(setDataReady);

    t1.join(); t2.join();
}
```